

PRACTICAL JOURNAL
in
Advanced Artificial Intelligence
&
Machine Learning
Submitted to

Laxman Devram Sonawane College, Kalyan (W) 421301

in partial fulfillment for the award of the degree of

Master of Science in Information Technology



(Affiliated to Mumbai University)

Submitted by
TEJAS ANIL KAPADE

Under the guidance of

Mrs. Sabina Ansari (Advanced Artificial Intelligence)

&

Dr. Priyanka Pawar (Machine Learning)

Department of Information Technology

Kalyan, Maharashtra

Academic Year 2025-26

PRACTICAL JOURNAL
in
Advance Artificial Intelligence
Submitted to

Laxman Devram Sonawane College, Kalyan (W) 421301

in partial fulfilment for the award of the degree of

Master of Science in Information Technology



(Affiliated to Mumbai University)

Submitted by

TEJAS ANIL KAPADE

Under the guidance of

Mrs. Sabina Ansari

Department of Information Technology
Kalyan, Maharashtra

Academic Year 2025-26



The Kalyan Wholesale Merchants Education Society's

Laxman Devram Sonawane College, Kalyan (W) 421301

**Department of Information Technology
Masters of Science – Part II**

Certificate

This is to certify that **Mr. Tejas Anil Kapade**, Seat number _____,
studying in Masters of Science in Information Technology Part II ,
Semester II has satisfactorily completed the practical of “Advanced Artificial
Intelligence ” as prescribed by University of Mumbai, during the academic year
2025-2026

Subject In-charge

Coordinator In-charge

External Examiner

College Seal

ADVANCED ARTIFICIAL INTELLIGENCE INDEX

SR No.	Practical	Sign
1	Implementing advanced deep learning algorithms such as convolutional neural networks (CNNs) or recurrent neural networks (RNNs) using Python libraries like TensorFlow or PyTorch.	
2	Building a natural language processing (NLP) model for sentiment analysis or text classification.	
3	Creating a chatbot using advanced techniques like transformer models	
4	Developing a recommendation system using collaborative filtering or deep learning approaches.	
5	Implementing a computer vision project, such as object detection or image segmentation	
6	Training a generative adversarial network (GAN) for generating realistic images	
7	Applying reinforcement learning algorithms to solve complex decision-making problems.	
8	Utilizing transfer learning to improve model performance on limited datasets.	
9	Building a deep learning model for time series forecasting or anomaly detection.	
10	Implementing a machine learning pipeline for automated feature engineering and model selection.	

11	Using advanced optimization techniques like evolutionary algorithms or Bayesian optimization for hyperparameter tuning	
12	Deploying a machine learning model in a production environment using containerization and cloud services.	
13	Use Python libraries such as GPT-2 or textgenrnn to train generative models on a corpus of text data and generate new text based on the patterns it has learned.	
14	Experiment with neural networks like GANs (Generative Adversarial Networks) using Python libraries like TensorFlow or PyTorch to generate new images based on a dataset of images.	

Practical 1

Aim : Implementing advanced deep learning algorithms such as convolutional neural networks (CNNs) or recurrent neural networks (RNNs) using Python libraries like TensorFlow or PyTorch

Code :

```
import tensorflow as tf
from tensorflow.keras import layers, models
import matplotlib.pyplot as plt

# Load and preprocess the CIFAR-10 dataset
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.cifar10.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0 # Normalize pixel values to [0, 1]

# Build the CNN model
model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'), layers.MaxPooling2D((2,
    2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.Flatten(), layers.Dense(64,
    activation='relu'),
    layers.Dense(10, activation='softmax') # 10 classes for CIFAR-10
])

# Compile the model
```

```
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
```

Train the model

```
history = model.fit(x_train, y_train, epochs=10, validation_data=(x_test, y_test))
```

```
# Evaluate the model test_loss, test_acc =
model.evaluate(x_test, y_test, verbose=2) print(f'\nTest
accuracy: {test_acc}')
```

```
# Plot training & validation accuracy values
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation
Accuracy') plt.title('Model accuracy') plt.ylabel('Accuracy')
plt.xlabel('Epoch') plt.legend() plt.show()
```

```
import torch import
torch.nn as nn import
torch.optim as optim import
torchvision
import torchvision.transforms as transforms import
matplotlib.pyplot as plt
```

Define transformations for the training and testing data


```
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)),
])
```

Load the CIFAR-10 dataset

trainset = torchvision

OUTPUT :

Test accuracy: 0.7017999887466431



Practical 2

Aim : Building a natural language processing (NLP) model for sentiment analysis or text classification.

Code :

```
!pip install tensorflow matplotlib
```

```
import tensorflow as tf
from tensorflow.keras import layers, models
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.datasets import imdb
import matplotlib.pyplot as plt
```

```
# Step 1: Load the IMDB dataset num_words = 10000 # Use the
top 10,000 words in the vocabulary
(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=num_words)
```

```
# Step 2: Explore the dataset
print(f'Number of training samples: {len(x_train)}')
print(f'Number of test samples: {len(x_test)}')
print(f'Sample review (tokenized): {x_train[0]}')
print(f'Label (0 = negative, 1 = positive): {y_train[0]}')
```

```
# Step 3: Decode a sample review word_index
= imdb.get_word_index()
reverse_word_index = {value: key for key, value in word_index.items()}
decoded_review = " ".join([reverse_word_index.get(i - 3, "?") for i in x_train[0]])
print(f'Decoded review: {decoded_review}')
```

```
# Step 4: Pad sequences maxlen = 200 # Limit
each review to 200 words x_train =
pad_sequences(x_train, maxlen=maxlen) x_test =
pad_sequences(x_test, maxlen=maxlen)
```

```
# Step 5: Define the model model
```

```
= models.Sequential([
    layers.Embedding(input_dim=num_words,                output_dim=32,
    input_length=maxlen), layers.LSTM(32), # Use an LSTM layer for capturing
    sequential dependencies layers.Dense(1, activation='sigmoid') # Output layer for
    binary classification
])
```

```
# Step 6: Compile the model model.compile(optimizer='adam',
    loss='binary_crossentropy',
    metrics=['accuracy'])
```

```
# Step 7: Display the model architecture
```

```
model.summary()
```

```
# Step 8: Train the model
```

```
history = model.fit(x_train, y_train, epochs=5, batch_size=64, validation_split=0.2)
```

```
# Step 9: Evaluate the model
```

```
test_loss, test_acc = model.evaluate(x_test, y_test, verbose=2) print(f'Test
Accuracy: {test_acc}')
```

```
# Step 10: Plot training history plt.figure(figsize=(12,
4))
```

```
# Accuracy plot plt.subplot(1,
2, 1)
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation
Accuracy') plt.xlabel('Epoch') plt.ylabel('Accuracy') plt.legend()
plt.title('Model Accuracy')
```

```
# Loss plot plt.subplot(1,
2, 2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Validation
Loss') plt.xlabel('Epoch') plt.ylabel('Loss') plt.legend()
plt.title('Model Loss') plt.show()
```



OUTPUT :

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb.npz>
17464789/17464789 — 0s 0us/step
Number of training samples: 25000
Number of test samples: 25000
Sample review (tokenized): [1, 14, 22, 16, 43, 530, 973, 1622, 1385, 65, 458, 4468, 66, 3941, 4, 173, 36, 256, 5, 25, 100, 43, 838, 112, 50, 670, 2, 9, 35, 480, 284, 5, 150, 4, 172,
Label (0 = negative, 1 = positive): 1
Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	?	0 (unbuilt)
lstm (LSTM)	?	0 (unbuilt)
dense (Dense)	?	0 (unbuilt)

Total params: 0 (0.00 B)
Trainable params: 0 (0.00 B)
Non-trainable params: 0 (0.00 B)
Epoch 1/5

Test Accuracy: 0.8611999750137329



Practical 3

Aim : Creating a chatbot using advanced techniques like transformer models.

Code :

```
!pip install transformers torch
```

```
from transformers import AutoModelForCausalLM, AutoTokenizer import  
torch
```

Step 1: Load Pre-trained Model and Tokenizer

```
print("Loading the DialoGPT model...") model_name =  
"microsoft/DialoGPT-medium"      tokenizer      =  
AutoTokenizer.from_pretrained(model_name)  
model = AutoModelForCausalLM.from_pretrained(model_name)
```

Step 2: Initialize Chat History

```
chat_history_ids = None step = 0
```

Step 3: Chat with the User

```
print("Chatbot is ready! Type 'exit' to end the chat.\n")
```

```
while True: # User input  
    user_input = input("You: ") if  
    user_input.lower() == 'exit':  
        print("Chatbot: Goodbye!")  
        break
```

```

# Encode the user input and add it to the chat history new_input_ids =
tokenizer.encode(user_input + tokenizer.eos_token, return_tensors='pt') chat_history_ids
= torch.cat([chat_history_ids, new_input_ids], dim=-1) if step > 0 else
new_input_ids

```

```

# Generate a response using the model response_ids =
model.generate(chat_history_ids, max_length=1000,
pad_token_id=tokenizer.eos_token_id) response =
tokenizer.decode(response_ids[:, chat_history_ids.shape[-1]:][0],
skip_special_tokens=True)

```

```

# Display the response
print(f"Chatbot: {response}")

```

```

step += 1

```

Output :

```

You: Hello
The attention mask is not set and cannot be inferred from input because pad token is same as eos token. As a consequence, you may observe unexpected behavior. Please pass your input
Chatbot: Hello! :D
You: How are you ?
Chatbot: I'm good, how are you?
You: What is your Name ?
Chatbot: I'm not sure.
You: What is your Qualification ?
Chatbot: I'm a human being.
You: What is your running speed ?
Chatbot: What is your favorite color?
You: Red & What is yours ?
Chatbot: I'm not sure.
You: exit
Chatbot: Goodbye!

```

Practical 4

Aim : Developing a recommendation system using collaborative filtering or deep learning approaches.

Code :

Step 1: Install Required Libraries

Run the following command to install the necessary libraries: pip
install tensorflow numpy pandas matplotlib

Step 2: Download the Dataset

Download the MovieLens 100K dataset from grouplens.org/datasets/movielens. Extract the dataset into a folder.

Alternatively, the code below assumes that the u.data file is in the ml-100k folder.

Step 3: Python Code for the Recommendation System

```
import pandas as pd
import numpy as np
import tensorflow as tf

from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt

# Step 1: Load and preprocess the dataset
file_path = "ml-100k/u.data"

column_names = ['user_id', 'item_id', 'rating', 'timestamp']
data = pd.read_csv(file_path, sep='\t', names=column_names)
```

```
# Normalize user IDs and item IDs to start from 0 data['user_id']
```

```
-= 1
```

```
data['item_id'] -= 1
```

```
# Extract the number of users and items
```

```
num_users = data['user_id'].max() + 1
```

```
num_items = data['item_id'].max() + 1
```

```
print(f'Number of users: {num_users}, Number of items: {num_items}')
```

```
# Step 2: Split data into training and testing
```

```
train_data, test_data = train_test_split(data, test_size=0.2, random_state=42)
```

```
# Step 3: Create TensorFlow datasets def
```

```
create_tf_dataset(df):
```

```
    users = tf.constant(df['user_id'].values, dtype=tf.int32)
```

```
    items = tf.constant(df['item_id'].values, dtype=tf.int32)
```

```
    ratings = tf.constant(df['rating'].values, dtype=tf.float32)
```

```
    return tf.data.Dataset.from_tensor_slices(((users, items), ratings)).shuffle(1024).batch(32)
```

```
train_dataset = create_tf_dataset(train_data) test_dataset
```

```
= create_tf_dataset(test_data)
```

```
# Step 4: Define the Recommendation Model class
```

```
MatrixFactorizationModel(tf.keras.Model):
```

```
    def __init__(self, num_users, num_items, embedding_dim=50):
```

```
        super().__init__() self.user_embedding = tf.keras.layers.Embedding(num_users,
```

```
        embedding_dim) self.item_embedding =
```

```
        tf.keras.layers.Embedding(num_items, embedding_dim)
```

```
    def call(self, inputs):
```



```

user_vector = self.user_embedding(inputs[0]) item_vector
= self.item_embedding(inputs[1])
dot_product = tf.reduce_sum(user_vector * item_vector, axis=1)
return dot_product

```

```

model = MatrixFactorizationModel(num_users, num_items)

```

```

#           Step           5:           Compile           the           model
model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=0.01),
loss='mse', metrics=['mae'])

```

```

# Step 6: Train the model history = model.fit(train_dataset,
validation_data=test_dataset, epochs=10)

```

```

# Step 7: Evaluate the model test_loss, test_mae =
model.evaluate(test_dataset) print(f'Test Loss:
{test_loss:.4f}, Test MAE: {test_mae:.4f}')

```

```

# Step 8: Plot the training history

```

```

plt.figure(figsize=(12, 4))

```

```

plt.subplot(1, 2, 1)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Validation
Loss') plt.xlabel('Epoch') plt.ylabel('Loss') plt.legend()
plt.title('Loss over Epochs')

```

```

plt.subplot(1, 2, 2)

```

```
plt.plot(history.history['mae'], label='Train MAE')
plt.plot(history.history['val_mae'], label='Validation
MAE') plt.xlabel('Epoch')
plt.ylabel('Mean Absolute Error')
plt.legend()
plt.title('MAE over Epochs')

plt.show()
```

Step 9: Make recommendations def

```
recommend(user_id, top_k=5):
```


```
    user_vector = tf.constant([user_id] * num_items, dtype=tf.int32)
    item_vector = tf.constant(list(range(num_items)), dtype=tf.int32)
    predictions = model.predict((user_vector, item_vector))
    top_items = np.argsort(-predictions)[:top_k]
    return top_items
```

```
user_id = 0 # Example user recommended_items
```

```
= recommend(user_id)
```

```
print(f'Recommended items for user {user_id}: {recommended_items}')
```

Output :



Practical 5

Aim : Implementing a computer vision project, such as object detection or image segmentation.

Code :

```
!pip install torch torchvision numpy matplotlib opencv-python ultralytics
```

```
import cv2 import numpy as
```

```
np import matplotlib.pyplot as
```

```
plt
```

```
from ultralytics import YOLO # YOLOv5 library from ultralytics
```

```
# Step 1: Load the YOLOv5 Model print("Loading
```

```
YOLOv5 model...")
```

```
model = YOLO("yolov5s.pt") # Use the small version of YOLOv5 pre-trained on COCO dataset
```

```
# Step 2: Load an Image for Object Detection
```

```
image_path = "/example.jpg" # Replace with your image file path image
```

```
= cv2.imread(image_path)
```

```
image_rgb = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
```

```
# Step 3: Perform Object Detection
```

```
print("Performing object detection...")
```

```
results = model.predict(image_rgb)
```

```
# Step 4: Visualize Results annotated_image = results[0].plot() # Annotated image with  
bounding boxes and labels plt.figure(figsize=(10, 10))
```

```
plt.imshow(cv2.cvtColor(annotated_image, cv2.COLOR_BGR2RGB))
```

```
plt.axis("off")
```

```
plt.title("Object Detection Results")
plt.show()
```

Step 5: Save the Annotated Image

```
output_path = "output.jpg"
cv2.imwrite(output_path,
annotated_image)
print(f'Annotated image saved to: {output_path}')
```

Step 6: Print Detected Objects

```
print("Detected objects:") for box in
results[0].boxes.data.tolist():
    x1, y1, x2, y2, conf, cls = box
    print(f'Class: {results[0].names[int(cls)]}, Confidence: {conf:.2f}, Coordinates: ({x1:.2f},
{y1:.2f}), ({x2:.2f}, {y2:.2f})')
```

Output :



Practical 6

Aim : Training a generative adversarial network (GAN) for generating realistic images

Code :

```
!pip install torch torchvision matplotlib numpy
```

```
import torch
import torch.nn as nn
import torch.optim as optim
from torchvision import datasets, transforms
from torch.utils.data import DataLoader
import matplotlib.pyplot as plt
```

Step 1: Define Generator and Discriminator class

Generator(nn.Module):

```
def _init_(self, noise_dim, img_dim):
    super(Generator, self)._init_()
    self.model = nn.Sequential(
        nn.Linear(noise_dim, 128),
        nn.ReLU(), nn.Linear(128,
        256), nn.ReLU(),
        nn.Linear(256, 512),
        nn.ReLU(), nn.Linear(512,
        img_dim), nn.Tanh(),
    )
def forward(self, x):
    return self.model(x)
```

```

class Discriminator(nn.Module):
    def __init__(self, img_dim):
        super(Discriminator, self).__init__()
        self.model = nn.Sequential(
            nn.Linear(img_dim, 512),
            nn.LeakyReLU(0.2),
            nn.Linear(512, 256),
            nn.LeakyReLU(0.2),
            nn.Linear(256, 1), nn.Sigmoid(),
        )

    def forward(self, x):
        return self.model(x)

```

Step 2: Define Constants and Hyperparameters

```

device = "cuda" if torch.cuda.is_available() else "cpu"
img_size = 28
img_dim = img_size * img_size
noise_dim = 100 batch_size =
64 epochs = 50
lr = 0.0002

```

Step 3: Prepare the MNIST Dataset

```

transform = transforms.Compose([transforms.ToTensor(), transforms.Normalize((0.5,),
(0.5,))])
dataset = datasets.MNIST(root="data", train=True, transform=transform, download=True)
dataloader = DataLoader(dataset, batch_size=batch_size, shuffle=True)

```

Step 4: Initialize Models, Loss, and Optimizers

```

generator = Generator(noise_dim,

```

```
img_dim).to(device) discriminator =
Discriminator(img_dim).to(device)
```

```
criterion = nn.BCELoss()
```

```
optimizer_g = optim.Adam(generator.parameters(), lr=lr) optimizer_d
= optim.Adam(discriminator.parameters(), lr=lr)
```

Step 5: Training Loop for

```
epoch in range(epochs):
```

```
for real_images, _ in dataloader: real_images =
real_images.view(-1, img_dim).to(device) batch_size =
real_images.size(0)
```

```
# Labels for real and fake images real_labels =
torch.ones(batch_size, 1).to(device) fake_labels =
torch.zeros(batch_size, 1).to(device)
```

Train Discriminator

```
noise = torch.randn(batch_size, noise_dim).to(device)
fake_images = generator(noise) real_preds =
discriminator(real_images) fake_preds =
discriminator(fake_images.detach()) loss_d_real =
criterion(real_preds, real_labels) loss_d_fake =
criterion(fake_preds, fake_labels) loss_d =
(loss_d_real + loss_d_fake) / 2
optimizer_d.zero_grad() loss_d.backward()
optimizer_d.step()
```

```
# Train Generator noise = torch.randn(batch_size,
noise_dim).to(device) fake_images =
generator(noise) fake_preds =
```

```

discriminator(fake_images)      loss_g      =
criterion(fake_preds,          real_labels)
optimizer_g.zero_grad() loss_g.backward()
optimizer_g.step()

```

Print progress

```
print(f'Epoch [{epoch+1}/{epochs}] | Loss D: {loss_d:.4f} | Loss G: {loss_g:.4f}')
```

Save generated samples every 10 epochs if

```

(epoch + 1) % 10 == 0:
    noise = torch.randn(16, noise_dim).to(device)
    generated_images = generator(noise).view(-1, 1, img_size, img_size).cpu().detach()
    plt.figure(figsize=(8, 8)) for i in range(16):
        plt.subplot(4, 4, i + 1)
        plt.imshow(generated_images[i].squeeze(),          cmap="gray")
        plt.axis("off")
    plt.tight_layout()
    plt.savefig(f'generated_images_epoch_{epoch+1}.png')
    plt.close()

```

Step 6: Save the Generator Model


```

torch.save(generator.state_dict(),
"gan_generator.pth") print("Generator model saved as
gan_generator.pth")


```


Output :


Generate after 10 Epoch



Generate after 20 Epoch




Generate 30 Epoch



Generate after 40 Epoch



Generate after 50 Epoch



Practical 7

Aim : Applying reinforcement learning algorithms to solve complex decision-making problems.

Code :

Step 1: Install Required Libraries

Run the following commands in your terminal to install the necessary libraries:

```
!pip install numpy gym matplotlib
```

Step 2: Python Code for Q-Learning

Save the following code as q_learning_cartpole.py.

```
import gym
import numpy as np
import matplotlib.pyplot as plt

# Step 1: Initialize Environment and Parameters
env = gym.make("CartPole-v1")
n_actions = env.action_space.n # Number of actions (2: left, right)
n_states = 20 # Discretize continuous state space
episodes = 500 # Number of episodes
learning_rate = 0.1 # Learning rate (alpha)
discount_factor = 0.99 # Discount factor (gamma)
epsilon = 1.0 # Exploration rate
epsilon_decay = 0.995 # Decay factor for epsilon
min_epsilon = 0.01 # Minimum epsilon

# Step 2: Helper Functions for Discretizing States
def discretize_state(state, state_bins):
```

```
return tuple(np.digitize(state[i], state_bins[i]) for i in range(len(state)))
```

```
def create_bins(n_states, env):
```

```
    state_bins = []
    for i in range(env.observation_space.shape[0]):
        low, high = env.observation_space.low[i], env.observation_space.high[i]
        bins = np.linspace(low, high, n_states - 1)
        state_bins.append(bins)
    return state_bins
```

```
# Step 3: Initialize Q-Table and State Bins state_bins
```

```
= create_bins(n_states, env)
```

```
q_table = np.zeros((n_states,) * len(env.observation_space.shape) + (n_actions,))
```

```
# Step 4: Training Loop
```

```
rewards = []
for episode in
```

```
range(episodes):
```

```
    state = discretize_state(env.reset()[0], state_bins)
    total_reward = 0
```

```
    done = False
```

```
    while not done:
```

```
        # Epsilon-greedy action selection if
```

```
        np.random.rand() < epsilon:
```

```
            action = np.random.choice(n_actions) # Explore else:
```

```
            action = np.argmax(q_table[state]) # Exploit
```

```
# Take action and observe results next_state_raw,
```

```
reward, done, _, _ = env.step(action)
next_state =
```

```
discretize_state(next_state_raw, state_bins)
```

```
total_reward += reward
```

```

# Q-Learning update q_table[state][action] += learning_rate * ( reward +
discount_factor * np.max(q_table[next_state]) - q_table[state][action]
)

```

```

state = next_state

```

```

# Decay epsilon

```

```

epsilon = max(min_epsilon, epsilon * epsilon_decay)

```

```

rewards.append(total_reward) print(f'Episode {episode + 1}/{episodes},

```

```

Reward: {total_reward}, Epsilon:

```

```

{epsilon:.3f}")

```

```

# Step 5: Plot Rewards plt.plot(rewards)

```

```

plt.title("Total Rewards Over Episodes")

```

```

plt.xlabel("Episode") plt.ylabel("Total

```

```

Reward") plt.show()

```

```

# Step 6: Save the Q-Table

```

```

np.save("q_table.npy", q_table)

```

```

print("Q-Table saved as 'q_table.npy'.")

```

```

# Step 7: Test the Trained Model

```

```

state = discretize_state(env.reset()[0], state_bins) done
= False

```

```

total_reward = 0 while

```

```

not done:

```

```

    action = np.argmax(q_table[state]) # Use trained Q-Table

```

```

    next_state_raw, reward, done, _, _ = env.step(action) state

```

```

    = discretize_state(next_state_raw, state_bins)

```

```

    total_reward += reward

```

```

    env.render()

```

```

env.close() print(f'Total reward during test:

```

```

{total_reward}")

```

Practical 8

Aim : Utilizing transfer learning to improve model performance on limited datasets.

Code :

Step 1: Install Required Libraries

Run the following command to install the necessary libraries:

```
!pip install torch torchvision matplotlib numpy
```

Step 2: Python Code for Transfer Learning

Save the following code as transfer_learning.py.

```
import torch
import torch.nn as nn
import torch.optim as optim
from torchvision import datasets, models, transforms
from torch.utils.data import DataLoader
import matplotlib.pyplot as plt

# Step 1: Set Device and Hyperparameters
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
num_classes = 10 # CIFAR-10 has 10 classes
batch_size = 32
epochs = 10
learning_rate = 0.001

# Step 2: Define Data Transformations
transform = transforms.Compose([
    transforms.Resize((224,
```

```
224)), # Resize to match ResNet input size
transforms.ToTensor(),
    transforms.Normalize(mean=[0.5, 0.5, 0.5], std=[0.5, 0.5, 0.5]),
])
```

```
# Step 3: Load CIFAR-10 Dataset train_dataset =
datasets.CIFAR10(root="data", train=True, transform=transform,
download=True) test_dataset = datasets.CIFAR10(root="data", train=False,
transform=transform, download=True)
```

```
train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True) test_loader
= DataLoader(test_dataset, batch_size=batch_size, shuffle=False)
```

```
# Step 4: Load Pre-trained ResNet18 Model
```

```
model = models.resnet18(pretrained=True)
for param in model.parameters():
    param.requires_grad = False # Freeze all layers
```

```
# Replace the final fully connected layer for CIFAR-10
```

```
model.fc = nn.Linear(model.fc.in_features, num_classes)
model = model.to(device)
```

```
# Step 5: Define Loss and Optimizer criterion
```

```
= nn.CrossEntropyLoss()
optimizer = optim.Adam(model.fc.parameters(), lr=learning_rate)
```

```
# Step 6: Training Loop def train(model,
```

```
loader, criterion, optimizer):
    model.train()
    running_loss = 0.0
    correct
```

= 0 total = 0 for inputs,

labels in loader:

```
inputs, labels = inputs.to(device), labels.to(device)
```

```
optimizer.zero_grad() outputs
```

```
= model(inputs) loss =
```

```
criterion(outputs, labels)
```

```
loss.backward()
```

```
optimizer.step()
```

```
running_loss += loss.item()
```

```
_, predicted =
```

```
outputs.max(1) total +=
```

```
labels.size(0)
```

```
correct += predicted.eq(labels).sum().item()
```

```
accuracy = 100. * correct / total
```

```
return running_loss / len(loader), accuracy
```

Step 7: Testing Loop def

test(model, loader, criterion):

```
model.eval()
```

```
running_loss = 0.0
```

```
correct = 0 total = 0
```

```
with
```

```
torch.no_grad():
```

```
for inputs, labels in loader:
```

```
inputs, labels = inputs.to(device), labels.to(device)
```



```

outputs = model(inputs) loss =
criterion(outputs, labels)
running_loss += loss.item()

_, predicted = outputs.max(1) total
+= labels.size(0)

correct += predicted.eq(labels).sum().item()

```

```

accuracy = 100. * correct / total
return running_loss / len(loader), accuracy

```

Step 8: Train and Evaluate

```

train_losses, test_losses = [], []
train_accuracies, test_accuracies = [], []

```

for epoch in range(epochs):

```

train_loss, train_acc = train(model, train_loader, criterion, optimizer) test_loss,
test_acc = test(model, test_loader, criterion)

```

```

train_losses.append(train_loss)
test_losses.append(test_loss)
train_accuracies.append(train_acc)
test_accuracies.append(test_acc)

```

```

print(f"Epoch {epoch+1}/{epochs}")
print(f"Train Loss: {train_loss:.4f}, Train Accuracy: {train_acc:.2f}%") print(f"Test
Loss: {test_loss:.4f}, Test Accuracy: {test_acc:.2f}%")

```

Step 9: Plot Training and Testing Metrics

```

plt.figure(figsize=(12, 5)) plt.subplot(1, 2, 1)

```

```
plt.plot(train_losses, label="Train Loss") plt.plot(test_losses,  
label="Test Loss") plt.legend() plt.title("Loss")
```

```
plt.subplot(1, 2, 2)  
plt.plot(train_accuracies, label="Train Accuracy")  
plt.plot(test_accuracies, label="Test Accuracy")  
plt.legend() plt.title("Accuracy") plt.show()
```

```
# Step 10: Save the Model  
torch.save(model.state_dict(),  
"resnet18_cifar10.pth") print("Model saved as  
'resnet18_cifar10.pth'.")
```

Practical 9

Aim : Building a deep learning model for time series forecasting or anomaly detection

Code :

```
!pip install numpy pandas matplotlib scikit-learn tensorflow
```

```
import numpy as np import
```

```
pandas as pd import
```

```
matplotlib.pyplot as plt
```

```
from sklearn.preprocessing import MinMaxScaler
```

```
from tensorflow.keras.models import Sequential
```

```
from tensorflow.keras.layers import LSTM, Dense
```

```
# Step 1: Load the Dataset url =
```

```
"https://raw.githubusercontent.com/jbrownlee/Datasets/master/airline-passengers.csv" data
```

```
= pd.read_csv(url, usecols=[1], header=0)
```

```
data = data.values.astype("float32") # Ensure the data is float
```

```
# Step 2: Visualize the Data plt.plot(data)
```

```
plt.title("Airline Passengers Over Time")
```

```
plt.xlabel("Time")
```

```
plt.ylabel("Passengers")
```

```
plt.show()
```

```
# Step 3: Normalize the Data scaler =
```

```
MinMaxScaler(feature_range=(0, 1))
```

```
scaled_data = scaler.fit_transform(data)
```

Step 4: Prepare the Data for LSTM

```
def create_dataset(dataset, look_back=1):
```

```
    X, y = [], []
    for i in range(len(dataset) -
```

```
        look_back):
```

```
        X.append(dataset[i:(i + look_back), 0])
```

```
        y.append(dataset[i + look_back, 0])
```

```
    return np.array(X), np.array(y)
```

```
look_back = 12 # Use 12 months (1 year) as input to predict the next value
```

```
X, y = create_dataset(scaled_data, look_back)
```

```
X = X.reshape((X.shape[0], X.shape[1], 1)) # Reshape for LSTM [samples, time_steps, features]
```

Step 5: Split Data into Training and Testing Sets

```
train_size = int(len(X) * 0.8)
```

```
X_train, X_test = X[:train_size], X[train_size:]
y_train, y_test = y[:train_size], y[train_size:]
```

Step 6: Build the LSTM Model

```
model = Sequential([
```

```
    LSTM(50, activation="relu", input_shape=(look_back, 1)), Dense(1)
```

```
])
```

```
model.compile(optimizer="adam", loss="mean_squared_error")
```

Step 7: Train the Model

```
history = model.fit(X_train, y_train, epochs=50,
```

```
                    batch_size=32, validation_data=(X_test, y_test), verbose=1)
```

Step 8: Evaluate the Model

```
loss, _ = model.evaluate(X_test, y_test, verbose=0)
print(f"Test
```

```
Loss: {loss:.4f}")
```

Step 9: Predict and Inverse Transform

```
y_pred = model.predict(X_test) y_pred =  
scaler.inverse_transform(y_pred)  
y_test_actual = scaler.inverse_transform(y_test.reshape(-1, 1))
```



Step 10: Plot Actual vs Predicted

```
plt.figure(figsize=(10, 5)) plt.plot(y_test_actual,  
label="Actual") plt.plot(y_pred,  
label="Predicted") plt.title("Actual vs Predicted  
Airline Passengers") plt.xlabel("Time")  
plt.ylabel("Passengers") plt.legend()  
plt.show()
```

Step 11: Save the Model

```
model.save("lstm_time_series.h5")  
print("Model saved as 'lstm_time_series.h5'.")
```

Output :



Practical 10

Aim : Implementing a machine learning pipeline for automated feature engineering and model selection.

Code :

Step 1: Install Required Libraries

Run the following command to install the required libraries

```
pip install pandas numpy scikit-learn
```

Step 2: Python Code for Machine Learning Pipeline

Save the following code as ml_pipeline.py.

```
import pandas as pd import numpy as np from
sklearn.model_selection import train_test_split, GridSearchCV from
sklearn.pipeline import Pipeline from sklearn.compose import
ColumnTransformer
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.feature_selection import SelectKBest, f_classif from
sklearn.ensemble import RandomForestClassifier from sklearn.svm
import SVC
from sklearn.metrics import classification_report, accuracy_score
```

Step 1: Load Dataset

Using the Titanic dataset for demonstration

```
url = "https://raw.githubusercontent.com/datasciencedojo/datasets/master/titanic.csv" data
= pd.read_csv(url)
```

Step 2: Basic Preprocessing

Drop unnecessary columns

```
data = data.drop(["PassengerId", "Name", "Ticket", "Cabin"], axis=1)
```

```
# Handle missing values data["Age"].fillna(data["Age"].median(),  
inplace=True)
```

```
data["Embarked"].fillna(data["Embarked"].mode()[0],  
inplace=True)
```

Separate features and target X

```
= data.drop("Survived", axis=1) y  
= data["Survived"]
```

Step 3: Split Data into Train and Test Sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Step 4: Define Preprocessing Steps

Numerical features: Scale values

```
numerical_features = ["Age", "Fare"]
```

```
numerical_transformer =
```

```
Pipeline(steps=[ ("scaler",
```

```
StandardScaler())
```

```
])
```

Categorical features: One-hot encode

```
categorical_features = ["Sex", "Embarked", "Pclass"]
```

```
categorical_transformer = Pipeline(steps=[
```

```
    ("onehot", OneHotEncoder(handle_unknown="ignore"))
```

```
])
```



```
# Combine preprocessors into a column transformer
preprocessor = ColumnTransformer(
transformers=[
    ("num", numerical_transformer, numerical_features),
    ("cat", categorical_transformer, categorical_features)
]
)
```

```
# Step 5: Define Feature Selection and Model Options feature_selection
= SelectKBest(score_func=f_classif)
```

```
# Define candidate models models
= {
    "RandomForest": RandomForestClassifier(random_state=42),
    "SVC": SVC(probability=True, random_state=42)
}
```

```
# Step 6: Create the Pipeline pipeline
= Pipeline(steps=[
    ("preprocessor", preprocessor),
    ("feature_selection", feature_selection),
    ("classifier", RandomForestClassifier())
])
```

```
# Step 7: Define Grid Search for Hyperparameter Tuning
```

```
param_grid = {
    "feature_selection_k": [5, 6, 7],
    "classifier": [models["RandomForest"], models["SVC"]],
    "classifier_____n_estimators": [100, 200] if "n_estimators" in
RandomForestClassifier().get_params() else [None],
```

```

    "classifier_C": [0.1, 1, 10] if "C" in SVC().get_params() else [None]
}

grid_search = GridSearchCV(pipeline, param_grid, cv=3, scoring="accuracy", verbose=2)
# Step 8: Train the Model
grid_search.fit(X_train, y_train)

# Step 9: Evaluate the Model best_model
= grid_search.best_estimator_ y_pred =
best_model.predict(X_test)

print("Best Parameters:", grid_search.best_params_) print("Accuracy
on Test Set:", accuracy_score(y_test, y_pred)) print("\nClassification
Report:\n", classification_report(y_test, y_pred))

```

Practical 11

Aim : Using advanced optimization techniques like evolutionary algorithms or Bayesian optimization for hyperparameter tuning.

Code :

```
!pip install numpy pandas scikit-learn scikit-optimize
```

```
import numpy as np
import pandas as pd

from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.ensemble import RandomForestClassifier
from skopt import BayesSearchCV

from sklearn.metrics import accuracy_score, classification_report
```

Step 1: Load the Dataset data

```
= load_iris()
X, y = data.data, data.target
```

Step 2: Split the Data into Training and Testing Sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Step 3: Define the Model

```
model = RandomForestClassifier(random_state=42)
```

Step 4: Define the Search Space for Hyperparameters

```
param_space = {
    "n_estimators": (10, 200),    # Number of trees in the forest
    "max_depth": (1, 20),        # Maximum depth of each tree
    "min_samples_split": (2, 10), # Minimum samples to split a node
    "min_samples_leaf": (1, 10),  # Minimum samples at each leaf
```

```
"max_features": ["sqrt", "log2", None] # Number of features considered for split
}
```

Step 5: Use Bayesian Optimization for Hyperparameter Tuning

```
optimizer = BayesSearchCV( estimator=model,
search_spaces=param_space, n_iter=30, # Number of iterations to
search cv=3,    # 3-fold cross-validation random_state=42, n_jobs=-
1
)
```

Step 6: Train the Optimized Model

```
print("Starting Bayesian Optimization...")
optimizer.fit(X_train, y_train)
```

Step 7: Evaluate the Best Model

```
best_model = optimizer.best_estimator_
y_pred = best_model.predict(X_test)
```

```
print("\nBest Parameters:", optimizer.best_params_) print("Accuracy
on Test Set:", accuracy_score(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

Optional: Save the Best Model import joblib

```
joblib.dump(best_model,
"optimized_rf_model.pkl") print("\nModel saved as
'optimized_rf_model.pkl'.")
```

Output :

```
Best Parameters: OrderedDict([('max_depth', 16), ('max_features', 'log2'), ('min_samples_leaf', 6), ('min_samples_split', 8), ('n_estimators', 182)])
Accuracy on Test Set: 1.0
```

```
Classification Report:
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	1.00	1.00	1.00	9
2	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

```
Model saved as 'optimized_rf_model.pkl'.
```

Practical 12

Aim : Deploying a machine learning model in a production environment using containerization and cloud services

Code :

Step 1: Install Required Libraries

```
!pip install scikit-learn pandas fastapi uvicorn
```

Step 2: Build the Machine Learning Model

Save the following Python script as train_model.py. This script trains the model and saves it.

```
import pandas as pd
from sklearn.datasets import load_iris
from sklearn.ensemble import RandomForestClassifier
import joblib

# Step 1: Load Dataset data
data = load_iris()
X, y = data.data, data.target

# Step 2: Train Model
model = RandomForestClassifier(random_state=42)
model.fit(X, y)

# Step 3: Save Model
joblib.dump(model, "iris_model.pkl")
print("Model saved as iris_model.pkl")

Run the script to save the trained model:
```

```
bash
Copy code python
train_model.py
```

Step 3: Create the API

Save the following Python script as app.py.

```
python Copy code from fastapi
import FastAPI from pydantic
import BaseModel import joblib
import numpy as np
```

Step 1: Load the trained model

```
model = joblib.load("iris_model.pkl")
```

Step 2: Initialize FastAPI

```
app = FastAPI()
```

Step 3: Define input schema class

```
IrisRequest(BaseModel):
```

```
    sepal_length: float
```

```
    sepal_width: float
```

```
    petal_length: float
```

```
    petal_width: float
```

Step 4: Define prediction endpoint

```
@app.post("/predict/") def
```

```
predict(iris: IrisRequest):
```

```
    features = np.array([[iris.sepal_length, iris.sepal_width, iris.petal_length,
iris.petal_width]])
```

```
    prediction = model.predict(features)
```

```
species = ["setosa", "versicolor", "virginica"]  
return {"prediction": species[prediction[0]]}
```

Run the API locally for testing: bash Copy code

```
uvicorn app:app --host 0.0.0.0 --port 8000
```

Test the API in your browser or with a tool like **Postman**:

□ URL: `http://127.0.0.1:8000/predict/` □ Example

Request Body:

json

Copy code

```
{  
  "sepal_length": 5.1,  
  "sepal_width": 3.5,  
  "petal_length": 1.4,  
  "petal_width": 0.2 }
```

Step 4: Create a Dockerfile Save

the following as Dockerfile.

dockerfile

Copy code

```
# Use an official Python runtime as the base image
```

```
FROM python:3.9-slim
```

```
# Set the working directory in the container
```

```
WORKDIR /app
```

```
# Copy the current directory contents into the container
```

```
COPY . /app
```

```
# Install required Python libraries
```

```
RUN pip install --no-cache-dir fastapi uvicorn scikit-learn joblib
```


Expose the API port

EXPOSE 8000

Command to run the application

CMD ["uvicorn", "app:app", "--host", "0.0.0.0", "--port", "8000"]

Step 5: Build and Run the Docker Container

1. Build the Docker image:

bash Copy

code

docker build -t iris-api .

2. Run the Docker container:

bash Copy code docker run -d -p

8000:8000 iris-api

3. Test the API:

- URL: `http://localhost:8000/predict/`
- Use the same JSON request as earlier.

Step 6: Deploy to a Cloud Service (Optional)

1. Prepare the Docker Image:

- Tag the image for a container registry (e.g., Docker Hub, AWS ECR, or GCP Artifact Registry):

bash Copy

code

docker tag iris-api <your_dockerhub_username>/iris-api docker

push <your_dockerhub_username>/iris-api

2. Deploy to AWS ECS (Example):

- Create an ECS cluster.
- Use the Docker image in a task definition.
- Deploy the task to the cluster.

3. Other Options:

- Use **AWS Lambda** with **API Gateway**.
 - Deploy on **Google Cloud Run** or **Azure App Service** for managed hosting.
-

Step 7: Dataset

□ Iris Dataset:

- Included with `sklearn.datasets` for demonstration purposes. ○

Automatically loaded in the `train_model.py` script.

Practical 13

Aim : Use Python libraries such as GPT-2 or textgenrnn to train generative models on a corpus of text data and generate new text based on the patterns it has learned.

Code :

Step 1: Install Required Libraries

Run the following commands to install necessary libraries:

```
!pip install transformers datasets torch
```

Step 2: Prepare a Text Dataset

For demonstration, we'll use the Tiny Shakespeare Corpus available via Hugging Face's datasets library. Alternatively, you can use your own dataset.

Step 3: Python Code for Training and Generating Text

Save the following code as train_gpt2.py.

```
import os

from datasets import load_dataset

from transformers import GPT2LMHeadModel, GPT2Tokenizer, Trainer, TrainingArguments

# Step 1: Load the Dataset
print("Loading dataset...")

dataset = load_dataset("tiny_shakespeare")

# Split into train and test sets

train_data = dataset["train"]

test_data = dataset["test"]
```

Step 2: Load Pre-trained GPT-2 Tokenizer and Model

```
print("Loading GPT-2 tokenizer and model...")
tokenizer = GPT2Tokenizer.from_pretrained("gpt2")
model = GPT2LMHeadModel.from_pretrained("gpt2")
```

Step 3: Tokenize the Dataset

```
def tokenize_function(examples):
    return tokenizer(examples["text"], truncation=True, padding="max_length",
max_length=512)

print("Tokenizing dataset...")
tokenized_train = train_data.map(tokenize_function, batched=True)
tokenized_test = test_data.map(tokenize_function, batched=True)
```

Step 4: Define Training Arguments

```
training_args = TrainingArguments(
    output_dir="./results",
    evaluation_strategy="epoch",
    learning_rate=5e-5,
    weight_decay=0.01,
    per_device_train_batch_size=4,
    num_train_epochs=3,
    save_total_limit=2,
    logging_dir="./logs",
    logging_steps=10,
)
```

Step 5: Initialize Trainer

```
trainer = GradientDescentTrainer(
```

```
model=model,  
args=training_args,  
train_dataset=tokenized_train,  
eval_dataset=tokenized_test,  
)
```

Step 6: Train the Model

```
print("Starting training...")  
trainer.train()
```

Step 7: Save the Fine-Tuned Model

```
model.save_pretrained("./fine_tuned_gpt2")  
tokenizer.save_pretrained("./fine_tuned_gpt2")  
print("Model saved to './fine_tuned_gpt2'.")
```

Step 8: Generate Text Using the Fine-Tuned Model

```
print("Generating new text...") model.eval()
```

```
input_text = "To be or not to be, that is the" inputs =  
tokenizer.encode(input_text, return_tensors="pt")  
outputs = model.generate(inputs, max_length=100, num_return_sequences=1, temperature=0.7)  
  
generated_text = tokenizer.decode(outputs[0], skip_special_tokens=True)  
print("\nGenerated Text:\n") print(generated_text)
```

Practical 14

Aim : Experiment with neural networks like GANs (Generative Adversarial Networks) using Python libraries like TensorFlow or PyTorch to generate new images based on a dataset of images.

Code :

```
!pip install torch torchvision matplotlib
```

```
import torch
import torch.nn as nn
import torch.optim as optim
from torchvision import datasets, transforms
from torch.utils.data import DataLoader
import matplotlib.pyplot as plt
import os
```

Step 1: Set Device Configuration

```
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(f"Using device: {device}")
```

Step 2: Define Generator class

```
Generator(nn.Module):
```

```
    def __init__(self, noise_dim, img_dim):
        super(Generator, self).__init__()
        self.gen = nn.Sequential(
            nn.Linear(noise_dim, 256),
            nn.ReLU(), nn.Linear(256,
            512), nn.ReLU(),
            nn.Linear(512, img_dim),
            nn.Tanh()
        )
```

```
def forward(self, x):  
    return self.gen(x)
```

Step 3: Define Discriminator class

```
Discriminator(nn.Module):  
    def __init__(self, img_dim):  
        super(Discriminator, self).__init__()  
        self.disc = nn.Sequential(  
            nn.Linear(img_dim, 512),  
            nn.LeakyReLU(0.2),  
            nn.Linear(512, 256),  
            nn.LeakyReLU(0.2),  
            nn.Linear(256, 1), nn.Sigmoid()  
        )  
  
    def forward(self, x): return  
        self.disc(x)
```

Step 4: Define Hyperparameters noise_dim

```
= 100  
  
img_dim = 28 * 28 # 28x28 images flattened  
batch_size = 64 lr = 0.0002  
epochs = 50
```

Step 5: Load Dataset

```
transform = transforms.Compose([transforms.ToTensor(), transforms.Normalize((0.5,),  
(0.5,))])  
  
dataset = datasets.MNIST(root="data", train=True, transform=transform, download=True)  
dataloader = DataLoader(dataset, batch_size=batch_size, shuffle=True)
```

Step 6: Initialize Models, Optimizers, and Loss Function

```
gen = Generator(noise_dim, img_dim).to(device) disc =  
Discriminator(img_dim).to(device) criterion = nn.BCELoss()  
opt_gen = optim.Adam(gen.parameters(), lr=lr) opt_disc  
= optim.Adam(disc.parameters(), lr=lr)
```

Step 7: Training Loop

```
print("Starting Training...")
```

```
for epoch in range(epochs):
```

```
    for batch_idx, (real, _) in enumerate(dataloader):
```

```
        real = real.view(-1, img_dim).to(device) batch_size  
        = real.size(0)
```

```
        # Train Discriminator
```

```
        noise = torch.randn(batch_size, noise_dim).to(device)
```

```
        fake = gen(noise) disc_real = disc(real).view(-1)
```

```
        disc_fake = disc(fake.detach()).view(-1)
```

```
        loss_disc = criterion(disc_real, torch.ones_like(disc_real)) + \ criterion(disc_fake,  
            torch.zeros_like(disc_fake))
```

```
        opt_disc.zero_grad()
```

```
        loss_disc.backward()
```

```
        opt_disc.step()
```

```
        # Train Generator
```

```
        output = disc(fake).view(-1)
```

```
        loss_gen = criterion(output,
```

```
            torch.ones_like(output)) opt_gen.zero_grad()
```

```
        loss_gen.backward() opt_gen.step()
```

```
print(f'Epoch [{epoch+1}/{epochs}] | Loss D: {loss_disc:.4f}, Loss G: {loss_gen:.4f}')
```



Save and Display Sample Images if

```
(epoch + 1) % 10 == 0:
    with torch.no_grad():
        fake_images = gen(torch.randn(16, noise_dim).to(device)).view(-1, 1, 28, 28)
    plt.figure(figsize=(4, 4))
    for i in range(16):
        plt.subplot(4, 4, i+1)
        plt.imshow(fake_images[i][0].cpu(),
                    cmap="gray")
        plt.axis("off")
    plt.tight_layout()
    os.makedirs("generated_images", exist_ok=True)
    plt.savefig(f"generated_images/epoch_{epoch+1}.png")
    plt.close()


print("Training Complete. Generated images are saved in 'generated_images' folder.")
```

Output :


Generate after 10 Epoch



Generate after 20 Epoch



Generate 30 Epoch



Generate after 40 Epoch



Generate after 50 Epoch



PRACTICAL JOURNAL

in

Machine Learning

Submitted to

Laxman Devram Sonawane College, Kalyan (W) 421301

in partial fulfilment for the award of the degree of

Master of Science in Information Technology



(Affiliated to Mumbai University)

Submitted by

TEJAS ANIL KAPADE

Under the guidance of

Dr. Priyanka Pawar

Department of Information Technology

Kalyan, Maharashtra

Academic Year 2025-26



The Kalyan Wholesale Merchants Education Society's

Laxman Devram Sonawane College, Kalyan (W) 421301

**Department of Information Technology
Masters of Science – Part II**

Certificate

This is to certify that **Mr. Tejas Anil Kapade**, Seat number _____,
studying in Masters of Science in Information Technology Part II ,
Semester III has satisfactorily completed the practical of “Advanced Artificial
Intelligence ” as prescribed by University of Mumbai, during the academic year
2025-2026

Subject In-charge

Coordinator In-charge

External Examiner

College Seal

MACHINE LEARNING

INDEX

Sr.no.	Practical	Date	Sign
1.	Data Pre-processing and Exploration a. Load a CSV dataset. Handle missing values, inconsistent formatting, and outliers. b. Load a dataset, calculate descriptive summary statistics, create visualizations using different graphs, and identify potential features and target variables Note: Explore Univariate and Bivariate graphs (Matplotlib) and Seaborn for visualization c. Create or Explore datasets to use all pre-processing routines like label encoding, scaling, and binarization.		
2.	Testing Hypothesis a. Implement and demonstrate the FIND-S algorithm for finding the most specific hypothesis based on a given set of training data samples. Read the training data from CSV file and generate the final specific hypothesis. (Create your dataset)		
3.	Linear Models a. Simple Linear Regression Fit a linear regression model on a dataset. Interpret coefficients, make predictions, and evaluate performance using metrics like R-squared and MSE b. Multiple Linear Regression Extend linear regression to multiple feature. Handle feature selection and potential multicollinearity c. Regularized Linear Models Implement Regression variants like LASSO and Ridge on any generated dataset		

4.	Discriminative Models a. Logistic Regression : Perform binary classification using logistic regression. Calculate accuracy, precision, recall, and understand the ROC curve." b. Implement and demonstrate k-nearest Neighbor algorithm. Read the training data from a .CSV file and build the model to classify a test sample. Print both correct and wrong predictions. c. Build a decision tree classifier or regressor. Control hyperparameters like tree depth to avoid overfitting. Visualize the tree. d. Implement a Support Vector Machine for any relevant dataset. e. Train a random forest ensemble. Experiment with the number of trees and feature sampling. Compare performance to a single decision tree. f. Implement a gradient boosting machine (e.g., XGBoost). Tune hyperparameters and explore feature importance.		
5.	Generative Models a. Implement and demonstrate the working of a Naive Bayesian classifier using a sample data set. Build the model to classify a test sample. b. Implement Hidden Markov Models using hmmlearn		
6.	Probabilistic Models a. Implement Bayesian Linear Regression to explore prior and posterior distribution. b. Implement Gaussian Mixture Models for density estimation and unsupervised clustering.		
7.	Model Evaluation and Hyperparameter Tuning a. Implement cross-validation techniques (k-fold, stratified, etc.) for robust model evaluation		

	b. Systematically explore combinations of hyperparameters to optimize model performance.(use grid and randomized search)		
8.	Bayesian Learning a. Implement Bayesian Learning using inferences		

Practical 1: Data Pre-processing and Exploration

1a. Load a CSV dataset. Handle missing values, inconsistent formatting, and outliers.

Code :

1. Import Libraries

```
# Import necessary libraries import
pandas as pd import numpy as np import
seaborn as sns
import matplotlib.pyplot as plt
```

2. Load the Dataset

```
# Load the Titanic dataset from a URL
url="https://raw.githubusercontent.com/datasciencedojo/datasets/master/titanic.csv" data =
pd.read_csv(url)
```

Display the first few rows

```
print(data.head())
```

3. Handle Missing Values # Check for missing values

```
print("Missing values in each column:")
print(data.isnull().sum())
```

```
# Fill missing values in 'Age' with the mean data['Age'].fillna(data['Age'].mean(),
inplace=True)
```

```
# Fill missing values in 'Embarked' with the most common value
data['Embarked'].fillna(data['Embarked'].mode()[0], inplace=True) #
```

```
Drop rows where 'Cabin' is missing (too many NaNs)
data.drop(columns=['Cabin'], inplace=True)
```

```
# Verify missing values are handled
```

```
print("\nAfter      handling      missing      values:")
print(data.isnull().sum())
```

4. Fix Inconsistent Formatting

```
# Fix inconsistent formatting in the 'Sex' column data['Sex']
= data['Sex'].str.lower().str.strip()
# Verify unique values print("\nUnique values in 'Sex'
column after formatting:") print(data['Sex'].unique())
```

5. Detect and Handle Outliers # Boxplot for the 'Fare' column

```
sns.boxplot(data['Fare'], color='skyblue') plt.title('Boxplot of
Fare') plt.show()
```

Detect outliers using the IQR method

```
Q1 = data['Fare'].quantile(0.25)
Q3 = data['Fare'].quantile(0.75)
IQR = Q3 - Q1 lower_bound =
Q1 - 1.5 * IQR upper_bound =
Q3 + 1.5 * IQR #
```


Capping outliers

```
data['Fare'] = np.where(data['Fare'] > upper_bound, upper_bound, np.where(data['Fare'] <
lower_bound, lower_bound, data['Fare'])) # Verify with an updated boxplot
sns.boxplot(data['Fare'], color='lightgreen') plt.title('Boxplot of Fare (After Handling
Outliers)') plt.show()
```

6. Save the Cleaned Dataset # Save the cleaned dataset

```
data.to_csv('cleaned_titanic.csv', index=False)
print("\nCleaned dataset saved as 'cleaned_titanic.csv'") .
```

Output :



1b. Load a dataset, calculate descriptive summary statistics, create visualizations using different graphs, and identify potential features and target variables Note:

Explore Univariate and Bivariate graphs (Matplotlib) and Seaborn for visualization Code :

1. Import Necessary Libraries # Import required libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

2. Load the Dataset

```
# Load the dataset from the URL url =
"https://raw.githubusercontent.com/mwaskom/seaborn-data/master/iris.csv" data
```

```
= pd.read_csv(url) # Display the first few rows print("First 5 rows of the  
dataset:") print(data.head())
```

3. Calculate Descriptive Summary Statistics # Dataset

```
information print("\nDataset Info:") print(data.info())  
  
# Summary statistics for numerical columns  
print("\nDescriptive Statistics for Numerical Columns:")  
print(data.describe())  
  
# Check unique values for categorical columns  
print("\nUnique values in 'species' column:")  
print(data['species'].value_counts())
```

4. Univariate Analysis

Histograms for numerical columns

```
data.hist(figsize=(10,8), color='skyblue',  
edgecolor='black') plt.suptitle("Histograms of Numerical  
Features") plt.show()
```

```
# Bar plot for 'species' column sns.countplot(x='species',  
data=data, palette='pastel')  
plt.title("Count of Each Species") plt.show()
```

5. Bivariate Analysis # Scatter plot for two features

```
plt.figure(figsize=(8, 6))  
  
plt.scatter(data['sepal_length'], data['sepal_width'], alpha=0.7, c='blue')  
plt.title("Sepal Length vs Sepal Width") plt.xlabel("Sepal Length")  
plt.ylabel("Sepal Width") plt.show()  
  
# Pairplot to visualize relationships between features  
sns.pairplot(data, hue='species', palette='husl',  
diag_kind='kde') plt.suptitle("Pairplot of Features by Species",  
y=1.02) plt.show()  
  
# Boxplot for petal_length across species sns.boxplot(x='species',  
y='petal_length', data=data, palette='Set3') plt.title("Boxplot of  
Petal Length by Species") plt.show()
```


6. Identify Potential Features and Target Variables # Separate

```
features and target features = data.drop(columns=['species'])  
  
# Drop the target column  
target = data['species'] #  
  
Target          variable  
  
print("\nFeatures:")  
print(features.head())  
print("\nTarget:")  
print(target.head())  
  
# Visualize target distribution  
sns.countplot(x=target,  
palette='viridis')      plt.title("Target  
Variable Distribution") plt.show()
```


7. Save the Cleaned and Processed Dataset # Save the dataset

```
data.to_csv('processed_iris.csv', index=False)  
  
print("\nProcessed dataset saved as 'processed_iris.csv'")
```


Output :





Count of Each Species



Sepal Length vs Sepal Width





1c. Create or Explore datasets to use all pre-processing routines like label encoding, scaling, and binarization.

Code :

1. Import Necessary Libraries # Import required libraries

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import LabelEncoder, MinMaxScaler, StandardScaler, Binarizer
```

2. Create or Load a Dataset # Create a sample dataset data

```
= pd.DataFrame({
    'Category': ['A', 'B', 'C', 'A', 'B', 'C'],
    # Categorical variable
    'Age': [23, 45, 31, 22, 35, 30],
    # Numerical variable
    'Income': [50000, 60000, 70000, 80000, 90000, 100000], #
    # Numerical variable 'Has_Car':
    'Has_Car': ['Yes', 'No', 'Yes', 'No', 'Yes', 'No']
    # Binary categorical variable })
```

Display the dataset

```
print("Sample Dataset:")
print(data)
```

3. Apply Pre-Processing Routines #

Label Encoding for 'Category' column

```
label_encoder = LabelEncoder()
data['Category_Encoded'] =
label_encoder.fit_transform(data['Category'])
```

Label Encoding for binary column 'Has_Car'

```
data['Has_Car_Encoded'] =
```

```

label_encoder.fit_transform(data['Has_Car']) print("\nAfter Label
Encoding:")

print(data)

# Min-Max Scaling for 'Income' min_max_scaler
= MinMaxScaler()

data['Income_MinMax'] = min_max_scaler.fit_transform(data[['Income']])

# Standard Scaling for 'Age' standard_scaler = StandardScaler()
data['Age_Standardized'] =
standard_scaler.fit_transform(data[['Age']]) print("\nAfter Scaling:")
print(data)

# Binarization for 'Income' with a threshold of 75,000

binarizer = Binarizer(threshold=75000)

data['Income_Binary'] =
binarizer.fit_transform(data[['Income']]) print("\nAfter
Binarization:")

print(data)

```

4. Save the Processed Dataset # Save the processed dataset

```

data.to_csv('processed_data.csv', index=False) print("\nProcessed
dataset saved as
'processed_data.csv'")

```

Output :



Sample Dataset:

	Category	Age	Income	Has_Car
0	A	23	50000	Yes
1	B	45	60000	No
2	C	31	70000	Yes
3	A	22	80000	No
4	B	35	90000	Yes
5	C	30	100000	No



After Label Encoding:

	Category	Age	Income	Has_Car	Category_Encoded	Has_Car_Encoded
0	A	23	50000	Yes	0	1
1	B	45	60000	No	1	0
2	C	31	70000	Yes	2	1
3	A	22	80000	No	0	0
4	B	35	90000	Yes	1	1
5	C	30	100000	No	2	0



After Scaling:

	Category	Age	Income	Has_Car	Category_Encoded	Has_Car_Encoded	\
0	A	23	50000	Yes	0	1	
1	B	45	60000	No	1	0	
2	C	31	70000	Yes	2	1	
3	A	22	80000	No	0	0	
4	B	35	90000	Yes	1	1	
5	C	30	100000	No	2	0	

	Income_MinMax	Age_Standardized
0	0.0	-1.035676
1	0.2	1.812434
2	0.4	0.000000
3	0.6	-1.165136
4	0.8	0.517838
5	1.0	-0.129460



After Binarization:

	Category	Age	Income	Has_Car	Category_Encoded	Has_Car_Encoded	\
0	A	23	50000	Yes	0	1	
1	B	45	60000	No	1	0	
2	C	31	70000	Yes	2	1	
3	A	22	80000	No	0	0	
4	B	35	90000	Yes	1	1	
5	C	30	100000	No	2	0	

	Income_MinMax	Age_Standardized	Income_Binary
0	0.0	-1.035676	0
1	0.2	1.812434	0
2	0.4	0.000000	0
3	0.6	-1.165136	1
4	0.8	0.517838	1
5	1.0	-0.129460	1

Practical 2 : Testing Hypothesis

Implement and demonstrate the FIND-S algorithm for finding the most specific hypothesis based on a given set of training data samples. Read the training data from a CSV file and generate the final specific hypothesis. (Create your dataset)

CODE :

```
import pandas as pd

# Step 1: Create the Dataset and Load It
data = {'Outlook': ['Sunny', 'Sunny', 'Overcast', 'Rainy', 'Rainy', 'Rainy', 'Overcast', 'Sunny', 'Sunny', 'Rainy'],
        'Temperature': ['Hot', 'Hot', 'Hot', 'Mild', 'Cool', 'Cool', 'Cool', 'Mild', 'Cool', 'Mild'],
        'Humidity': ['High', 'High', 'High', 'High', 'Normal', 'Normal', 'Normal', 'High', 'Normal', 'Normal'],
        'Wind': ['Weak', 'Strong', 'Weak', 'Weak', 'Weak', 'Strong', 'Strong', 'Weak', 'Weak', 'Weak'],
        'PlayTennis': ['No', 'No', 'Yes', 'Yes', 'Yes', 'No', 'Yes', 'No', 'Yes', 'Yes']}

# Load dataset into a pandas DataFrame df
df = pd.DataFrame(data)

# Step 2: Implementing the FIND-S Algorithm
def find_s_algorithm(data):
    # Get the positive examples (PlayTennis = 'Yes')
    positive_examples = data[data['PlayTennis'] == 'Yes']

    # Initialize hypothesis with the first positive example (most specific)
    hypothesis = positive_examples.iloc[0].drop('PlayTennis')

    # Loop through the rest of the positive examples and generalize the hypothesis
    for index, row in positive_examples.iterrows():
        for feature in hypothesis.index:
            if hypothesis[feature] != row[feature]:
                hypothesis[feature] = '?'
    return hypothesis
```

Apply FIND-S to the dataset

```
hypothesis = find_s_algorithm(df) #
```

Display the final specific hypothesis

```
print("The most specific hypothesis is:")
```

```
print(hypothesis)
```

Output :

```
Dataset:
  Sky Temperature Humidity Wind Water Forecast Condition
0 Sunny      Warm   Normal Strong  Warm     Same        Yes
1 Sunny      Cold    High  Strong  Warm     Same        No
2 Rainy      Warm    High   Weak   Cool    Change       No
3 Sunny      Warm   Normal Strong  Warm     Same        Yes
4 Rainy      Cold   Normal   Weak   Cool    Change       No
```

```
Loaded Dataset:
  Sky Temperature Humidity Wind Water Forecast Condition
0 Sunny      Warm   Normal Strong  Warm     Same        Yes
1 Sunny      Cold    High  Strong  Warm     Same        No
2 Rainy      Warm    High   Weak   Cool    Change       No
3 Sunny      Warm   Normal Strong  Warm     Same        Yes
4 Rainy      Cold   Normal   Weak   Cool    Change       No
```

```
Final Specific Hypothesis:
['Sunny', 'Warm', 'Normal', 'Strong', 'Warm', 'Same']
```

Practical 3 : Linear Models

3a. Simple Linear Regression

Fit a linear regression model on a dataset. Interpret coefficients, make predictions, and evaluate performance using metrics like R-squared and MSE

Code :

Step 1: Import Libraries # Import required libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression

from sklearn.metrics import mean_squared_error, r2_score
```

Step 2: Create a Dataset and Save as CSV

Create a sample dataset data

```
= {
    'House_Size': [750, 800, 850, 900, 1000, 1100, 1200, 1300, 1400, 1500],
    'Price': [150000, 160000, 165000, 170000, 180000, 190000, 200000, 210000, 220000, 230000]
}
```

Convert the dataset into a DataFrame

```
df = pd.DataFrame(data)
```

Save to CSV file

```
df.to_csv('house_prices.csv',
index=False)
```

Display the dataset

```
print("Dataset:")
print(df)
```

Step 3: Load the Dataset

```
# Load the dataset dataset =  
pd.read_csv('house_prices.csv')  
  
# Display the first few rows  
print("\nLoaded  
Dataset:")  
print(dataset.head())
```

Step 4: Split the Dataset into Training and Test Sets

Features and target variable

```
X = dataset[['House_Size']] # Feature: House size y  
= dataset['Price']          # Target: Price
```

Split data into training and testing sets (80% train, 20% test)

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)  
print("\nTraining and Testing Data Sizes:") print("Training Data Size:",  
X_train.shape[0]) print("Testing Data Size:", X_test.shape[0])
```

Step 5: Fit a Linear Regression Model

Initialize and fit the linear regression model

```
model = LinearRegression() model.fit(X_train,  
y_train) # Display the coefficients  
print("\nModel Coefficients:") print("Slope  
(m):", model.coef_[0]) print("Intercept (b):",  
model.intercept_)
```

Step 6: Make Predictions

Predict on the test set

```
y_pred = model.predict(X_test) #  
Display predictions  
print("\nPredictions on Test Data:")  
print("Actual Prices:",  
y_test.values) print("Predicted  
Prices:", y_pred)
```

Step 7: Evaluate the Model # Calculate evaluation metrics

```
mse = mean_squared_error(y_test, y_pred) r2 = r2_score(y_test, y_pred) #
```

Display metrics

```
print("\nModel Performance Metrics:")
print("Mean Squared Error (MSE):", mse)
print("R-squared ( $R^2$ ):", r2)
```

Step 8: Visualize the Results # Scatter plot of the training


```
data plt.scatter(X_train, y_train, color='blue', label='Training
Data')
```

```
# Plot the regression line plt.plot(X_train, model.predict(X_train),
color='red', label='Regression Line')
```

Scatter plot of the test data

```
plt.scatter(X_test, y_test, color='green', label='Test Data')
plt.title("Simple Linear Regression") plt.xlabel("House
Size (sq ft)") plt.ylabel("Price ($)") plt.legend()
plt.show()
```

Output :



3b. Multiple Linear Regression :

Extend linear regression to multiple feature. Handle feature selection and potential multicollinearity

Code :

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.preprocessing import LabelEncoder

# Import LabelEncoder

from sklearn.impute import SimpleImputer

# Load dataset from google.colab import files
uploaded = files.upload()

Upload your CSV file

# Read the CSV file data

data = pd.read_csv(list(uploaded.keys())[0])

# Display the first few rows print(data.head())

# Check for null values and basic statistics
print(data.info())
print(data.describe())

# Define a function to calculate VIF
def calculate_vif(df):

# Select only numeric features for VIF calculation
numeric_df = df.select_dtypes(include=[np.number])

# Drop rows with infinite or missing values

numeric_df = numeric_df.replace([np.inf, -np.inf], np.nan).dropna()
vif_data = pd.DataFrame()
vif_data["feature"] = numeric_df.columns
vif_data["VIF"] = [variance_inflation_factor(numeric_df.values,
i) for i in range(numeric_df.shape[1])]
return vif_data
```

Selecting features and target variable

```
X = data.drop("Survived", axis=1)
# Changed 'y' to 'Survived' y = data["Survived"]
```

Handle categorical features (e.g., using Label Encoding)

```
for col in X.select_dtypes(include=['object']).columns: le =
LabelEncoder()
    X[col] = le.fit_transform(X[col])
```

```
# Impute missing values using the mean (you can choose other strategies) imputer
= SimpleImputer(strategy='mean')
```

Create an imputer instance

```
X = pd.DataFrame(imputer.fit_transform(X), columns=X.columns)
```

Impute and update X

```
# Calculate VIF for initial features print("VIF
before      handling      multicollinearity:")
print(calculate_vif(X)) # Call the modified function
```

Drop features based on VIF analysis (example: drop 'X1' if VIF is high)

Check if the column exists before dropping if

```
'X1' in X.columns:
```

```
    X = X.drop("X1", axis=1) # Replace 'X1' with the actual high VIF feature name else:
    print("Column 'X1' not found in the DataFrame.")
```

Recalculate VIF

```
print("VIF after handling multicollinearity:")
print(calculate_vif(X))
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Initialize and fit the model

```
model = LinearRegression()
model.fit(X_train, y_train)
```

Get coefficients and intercept

```
print("Coefficients:", model.coef_)
print("Intercept:",
model.intercept_)
```

Predictions

```
y_pred = model.predict(X_test)
```

Evaluation metrics

```
rmse = np.sqrt(mean_squared_error(y_test, y_pred)) r2 = r2_score(y_test, y_pred)
print(f"RMSE: {rmse}") print(f"R^2: {r2}")
from sklearn.feature_selection import RFE
```

Recursive Feature Elimination

```
rfe = RFE(estimator=LinearRegression(), n_features_to_select=5)
```

```
# Adjust features
```

```
rfe.fit(X_train, y_train)
```

Selected features

```
print("Selected Features:", X.columns[rfe.support_])
```

Scatter plot of actual vs predicted values

```
plt.scatter(y_test, y_pred)
```

```
plt.xlabel("Actual") plt.ylabel("Predicted")
```


```
plt.title("Actual vs Predicted") plt.show()
```


Residuals

```
residuals = y_test - y_pred sns.histplot(residuals, kde=True)
```

```
plt.title("Residuals Distribution") plt.show()
```

Output :





3c. Regularized Linear Models :

Implement Regression variants like LASSO and Ridge on any generated dataset

Code :

1. Set Up the Environment

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.linear_model import Ridge, Lasso, ElasticNet
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.datasets import make_regression # Set random seed for reproducibility

np.random.seed(42)
```

2. Generate a Synthetic Dataset

```
# Generate synthetic data
X, y = make_regression(n_samples=1000,

# Number of samples
n_features=10, #
Number of features
noise=15,

# Add some noise
random_state=42
)

# Convert to DataFrame for exploration
data = pd.DataFrame(X, columns=[f'X{i}'
for i in range(1, 11)])
data["y"] = y

# Display the first few rows
```

```
print(data.head())
```

3. Split the Dataset

Split data into training and testing sets

```
X_train, X_test, y_train, y_test = train_test_split(data.drop("y", axis=1),
```

```
# Features data["y"],
```

```
# Target variable test_size=0.2,
```

```
# 20% for testing random_state=42
```

```
)
```

4. Train and Evaluate Ridge Regression

Initialize Ridge Regression with a regularization parameter (alpha)

```
ridge = Ridge(alpha=1.0) # Train the model ridge.fit(X_train, y_train) #
```

Predictions

```
ridge_pred = ridge.predict(X_test) #
```

Evaluate Ridge Regression

```
ridge_rmse = np.sqrt(mean_squared_error(y_test, ridge_pred))
```

```
ridge_r2 = r2_score(y_test, ridge_pred) print(f'Ridge RMSE:
```

```
{ridge_rmse}") print(f'Ridge R^2: {ridge_r2}")
```

5. Train and Evaluate Lasso Regression # Initialize Lasso Regression

```
lasso = Lasso(alpha=0.1)
```

```
# Train the model lasso.fit(X_train,
```

```
y_train) # Predictions
```

```
lasso_pred = lasso.predict(X_test)
```

```
# Evaluate Lasso Regression lasso_rmse =
```

```
np.sqrt(mean_squared_error(y_test, lasso_pred)) lasso_r2 =
```

```
r2_score(y_test, lasso_pred) print(f'Lasso RMSE:
```

```
{lasso_rmse}") print(f'Lasso R^2: {lasso_r2}") # Features
```

shrunk to

```
zero print("Lasso Coefficients:", lasso.coef_)
```

6. Train and Evaluate ElasticNet Regression

Initialize ElasticNet

```
elastic_net = ElasticNet(alpha=0.1, l1_ratio=0.5) # l1_ratio balances L1 and L2 penalties
```

```
# Train the model elastic_net.fit(X_train,  
y_train)
```

Predictions

```
elastic_net_pred = elastic_net.predict(X_test)
```

Evaluate

```
ElasticNet Regression elastic_net_rmse = np.sqrt(mean_squared_error(y_test,  
elastic_net_pred)) elastic_net_r2 = r2_score(y_test, elastic_net_pred)  
print(f'ElasticNet RMSE: {elastic_net_rmse}') print(f'ElasticNet R^2:  
{elastic_net_r2}')
```

7. Compare Results


Collect metrics metrics

```
= pd.DataFrame({  
    "Model": ["Ridge", "Lasso", "ElasticNet"],  
    "RMSE": [ridge_rmse, lasso_rmse, elastic_net_rmse],  
    "R^2": [ridge_r2, lasso_r2, elastic_net_r2]  
})
```

```
print(metrics)
```

```
#          Plot      RMSE      comparison  
sns.barplot(data=metrics,          x="Model",  
y="RMSE") plt.title("Model RMSE Comparison")  
plt.show()
```

Output :



Practical 4 : Discriminative Models

4a. Logistic Regression :

Perform binary classification using logistic regression. Calculate accuracy, precision, recall, and understand the ROC curve."

Code :

Step 1: Import Required Libraries # Import

```
necessary libraries import numpy as np import
pandas as pd from sklearn.model_selection import
train_test_split from sklearn.linear_model import
LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score, roc_curve, auc
import matplotlib.pyplot as plt
```

Step 2: Prepare the Dataset from

```
sklearn.datasets import make_classification
```

Create a synthetic dataset

```
X, y = make_classification(n_samples=1000, n_features=10, n_classes=2, random_state=42)
```

Split data into training and testing sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

Step 3: Train the Logistic Regression Model

Initialize the logistic regression model logreg

```
= LogisticRegression()
```

```
# Train the model on the training data logreg.fit(X_train,
y_train)
```

Step 4: Make Predictions

Predict labels for the test set y_pred


```
= logreg.predict(X_test)
```

```
# Predict probabilities for the ROC curve
y_prob = logreg.predict_proba(X_test)[:, 1]
```

Step 5: Evaluate the Model # Calculate

```
metrics accuracy = accuracy_score(y_test,
y_pred) precision = precision_score(y_test,
y_pred) recall = recall_score(y_test,
y_pred) print(f'Accuracy:
{accuracy:.2f}')
```

Output :



4b .Implement and demonstrate k-nearest Neighbor algorithm. Read the training data from a .CSV file and build the model to classify a test sample. Print both correct and wrong predictions.

Code :

Step 1: Import Required Libraries # Import necessary libraries

```
import pandas as pd
import numpy as np

from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score

from google.colab import files
```

Step 2: Create or Upload the CSV File

Check if the user wants to create a dataset or upload one

```
print("Do you have a CSV file to upload? (yes/no)")
response = input().lower()
if response == "yes":
```

```
    uploaded = files.upload()
    filename = list(uploaded.keys())[0]
else:
```

Create a synthetic dataset from

```
sklearn.datasets import make_classification
```

Generate synthetic data

```
X,y=make_classification(n_samples=200,n_features=5, n_classes=2, random_state=42)
```

Combine features and target into a single DataFrame

```
data = pd.DataFrame(X, columns=[f'Feature_{i}' for i in
range(X.shape[1])])
data['Target'] = y
```

Save the dataset to a CSV file filename =

```
"synthetic_data.csv"
data.to_csv(filename,
index=False)
print(f'Synthetic dataset saved as {filename}.')
Step 3: Load the CSV File into
```

a DataFrame

Load the dataset into a DataFrame

```
data = pd.read_csv(filename)
```

Display the first few rows of the dataset

```
print("Loaded Dataset:")
print(data.head())
```

Step 4: Preprocess the Data

Separate features (X) and labels (y)

```
X = data.iloc[:, :-1].values # All columns except the last one
y = data.iloc[:, -1].values # Last column as the target
```

Split the dataset into training and testing sets (80% train, 20% test)

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Step 5: Train the k-NN Model #

Initialize the k-NN model with k=3 knn

```
= KNeighborsClassifier(n_neighbors=3) #
```

Train the model on the training data

```
knn.fit(X_train, y_train)
```

Step 6: Predict Test Samples #

Predict the labels for the test set

```
y_pred = knn.predict(X_test)
```

Step 7: Evaluate and Print Predictions #

Calculate and display the accuracy

```
accuracy = accuracy_score(y_test, y_pred)
print(f'\nModel Accuracy:
{accuracy:.2f}\n')
```

Display correct and incorrect predictions

```
print("Correct Predictions:")
for i in range(len(y_test)):
    if y_pred[i] == y_test[i]:
        print(f'Sample {i}: Predicted={y_pred[i]}, Actual={y_test[i]}')
print("\nIncorrect Predictions:")
for i in range(len(y_test)):
    if y_pred[i] != y_test[i]:
        print(f'Sample {i}: Predicted={y_pred[i]}, Actual={y_test[i]}')
```

Output :

```
[ ] Model Accuracy: 0.88
Correct Predictions:
Sample 0: Predicted=0, Actual=0
Sample 1: Predicted=1, Actual=1
Sample 2: Predicted=1, Actual=1
Sample 3: Predicted=0, Actual=0
Sample 4: Predicted=1, Actual=1
Sample 5: Predicted=1, Actual=1
Sample 6: Predicted=0, Actual=0
Sample 7: Predicted=0, Actual=0
Sample 9: Predicted=1, Actual=1
Sample 10: Predicted=1, Actual=1
Sample 11: Predicted=1, Actual=1
Sample 12: Predicted=0, Actual=0
Sample 13: Predicted=0, Actual=0
Sample 14: Predicted=0, Actual=0
Sample 15: Predicted=0, Actual=0
Sample 16: Predicted=0, Actual=0
Sample 17: Predicted=1, Actual=1
Sample 18: Predicted=1, Actual=1
Sample 19: Predicted=0, Actual=0
Sample 20: Predicted=0, Actual=0
Sample 22: Predicted=1, Actual=1
Sample 23: Predicted=1, Actual=1
Sample 24: Predicted=1, Actual=1
Sample 25: Predicted=1, Actual=1
Sample 26: Predicted=1, Actual=1
Sample 27: Predicted=0, Actual=0
Sample 28: Predicted=0, Actual=0
Sample 30: Predicted=1, Actual=1
Sample 31: Predicted=1, Actual=1
Sample 32: Predicted=1, Actual=1
Sample 34: Predicted=0, Actual=0
Sample 35: Predicted=1, Actual=1
Sample 36: Predicted=1, Actual=1
Sample 38: Predicted=1, Actual=1
Sample 39: Predicted=1, Actual=1
```

```
Incorrect Predictions:
Sample 8: Predicted=1, Actual=0
Sample 21: Predicted=1, Actual=0
Sample 29: Predicted=0, Actual=1
Sample 33: Predicted=1, Actual=0
Sample 37: Predicted=1, Actual=0
```

4c. Build a decision tree classifier or regressor. Control hyperparameters like tree depth to avoid overfitting. Visualize the tree.

Code :

Step 1: Import Required Libraries

Import necessary libraries

```
import pandas as pd
import numpy
```

```
as np
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.tree import DecisionTreeClassifier, DecisionTreeRegressor, plot_tree
```

```
from sklearn.metrics import accuracy_score, mean_squared_error
```

```
import matplotlib.pyplot as plt
```

```
from google.colab import files
```

Step 2: Create or Upload the CSV File

Check if the user wants to upload a file or generate one

```
print("Do you have a CSV file to upload? (yes/no)") response  
= input().lower() if response == "yes":
```

Upload the CSV file uploaded =

```
files.upload() filename =
```

```
list(uploaded.keys())[0]
```

```
else:
```

Generate synthetic data (classification or regression)

```
from sklearn.datasets import make_classification, make_regression  
print("Choose a task: (1) Classification (2) Regression")
```

```
task = int(input()) if task == 1:
```

Generate synthetic classification data

```
X, y = make_classification(n_samples=200, n_features=5, random_state=42) task_type  
= "classification"
```

```
else:
```

Generate synthetic regression data

```
X, y = make_regression(n_samples=200, n_features=5, random_state=42) task_type  
= "regression"
```

Combine features and target into a single DataFrame

```
data = pd.DataFrame(X, columns=[f"Feature_{i}" for i in  
range(X.shape[1])]) data['Target'] = y
```

Save the dataset to a CSV file

```
filename = "synthetic_data.csv"
```

```
data.to_csv(filename, index=False)
```

```
print(f'Synthetic {task_type} dataset saved as {filename}.')
```

Step 3: Load the Dataset

Load the dataset data

```
=pd.read_csv(filename)
```

Display the first few rows of the dataset

```
print("Dataset Preview:") print(data.head())
```

Step 4: Preprocess the Data

Separate features and target

```
X = data.iloc[:, :-1].values # All columns except the last one
```

```
y = data.iloc[:, -1].values # Last column as the target
```

Split data into training and testing sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Step 5: Build the Decision Tree

Define the tree depth to avoid overfitting max_depth

```
= 3
```

Initialize the model

```
if task_type == "classification": model =
```

```
DecisionTreeClassifier(max_depth=max_depth, random_state=42) else:
```

```
model = DecisionTreeRegressor(max_depth=max_depth, random_state=42)
```

Train the model

```
model.fit(X_train, y_train)
```

Step 6: Make Predictions

Predict on the test set y_pred =

```
model.predict(X_test) #
```

Evaluate the model if task_type
== "classification":

```
accuracy = accuracy_score(y_test, y_pred)
```

```
print(f"Accuracy: {accuracy:.2f}")
```


else:

```
mse = mean_squared_error(y_test, y_pred)
```

```
print(f"Mean Squared Error: {mse:.2f}")
```

Step 7: Visualize the Tree # Visualize the decision tree `plt.figure(figsize=(12, 8))`
`plot_tree(model,` `feature_names=data.columns[:-1],`
`class_names=str(np.unique(y)) if task_type == "classification" else None,`
`filled=True)` `plt.title("Decision Tree Visualization")` `plt.show()`

Output :



4d. Implement a Support Vector Machine for any relevant dataset.

Code:

Step 1: Import Required Libraries # Import

necessary libraries `import pandas as pd` `import`

`numpy as np` `from sklearn.model_selection` `import`

`train_test_split` `from sklearn.svm` `import SVC`

`from sklearn.metrics` `import accuracy_score, classification_report` `from`
`google.colab` `import files`

Step 2: Create or Upload a Dataset

Check if the user wants to upload a file or generate one

```
print("Do you have a CSV file to upload? (yes/no)") response  
= input().lower() if response == "yes": # Upload the CSV  
file    uploaded    =    files.upload()    filename    =  
list(uploaded.keys())[0] else:
```

Generate synthetic classification data from

```
sklearn.datasets import make_classification
```

```
X, y = make_classification(n_samples=200, n_features=5, n_classes=2, random_state=42)
```

Combine features and target into a DataFrame

```
data = pd.DataFrame(X, columns=[f'Feature_{i}'  
for i in range(X.shape[1])]) data['Target'] = y
```

#Save the synthetic dataset to a CSV file

```
filename="synthetic_data.csv"
```

```
data.to_csv(filename,index=False)
```

```
print(f'Synthetic dataset saved as {filename}.')
```

Step 3: Load the Dataset

Load the dataset into a DataFrame data

```
= pd.read_csv(filename)
```

Display the first few rows of the dataset

```
print("Dataset Preview:") print(data.head())
```

Step 4: Preprocess the Data

Separate features (X) and target (y)

```
X = data.iloc[:, :-1].values # All columns except the last one
```

```
y = data.iloc[:, -1].values # Last column as the target
```

Split the dataset into training (80%) and testing (20%) sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Step 5: Train the Support Vector Machine

```
# Initialize the SVM model (use RBF kernel as default) svm_model =  
SVC(kernel='rbf', C=1.0, gamma='scale', random_state=42)  
  
# Train the SVM model on the training data  
svm_model.fit(X_train, y_train)
```

Step 6: Make Predictions

```
# Predict the labels for the test set  
y_pred = svm_model.predict(X_test)
```

Step 7: Evaluate the Model # Calculate

```
and print the accuracy accuracy =  
accuracy_score(y_test, y_pred)  
print(f'Model Accuracy: {accuracy:.2f}')  
  
# Print a detailed classification report  
print("\nClassification Report:")  
print(classification_report(y_test, y_pred))
```


Step 8: Visualize the Decision Boundary (Optional for 2D

```
Data) import matplotlib.pyplot as plt  
  
# Generate 2D synthetic data from  
sklearn.datasets import make_blobs  
  
X, y = make_blobs(n_samples=100, centers=2, random_state=42, cluster_std=1.5)  
  
# Fit the SVM on this data svm_model.fit(X,  
y)  
  
#Plot the decision boundary plt.figure(figsize=(8,  
6))  
  
plt.scatter(X[:, 0], X[:, 1], c=y, cmap='coolwarm', edgecolor='k')  
# Create a grid to evaluate the model xx, yy = np.meshgrid(np.linspace(X[:,  
0].min(), X[:, 0].max(), 100), np.linspace(X[:, 1].min(), X[:, 1].max(), 100))  
Z = svm_model.decision_function(np.c_[xx.ravel(), yy.ravel()])  
  
Z = Z.reshape(xx.shape)
```

Plot the decision boundary and margins

```
plt.contour(xx, yy, Z, levels=[-1, 0, 1], linestyles=['--', '-', '--'], colors='k')  
plt.title("SVM Decision Boundary") plt.xlabel("Feature 1")  
plt.ylabel("Feature 2") plt.show()
```

Output :



4e. Train a random forest ensemble. Experiment with the number of trees and feature sampling. Compare performance to a single decision tree.

Code :

Step 1: Import Required Libraries # Import necessary libraries

```
import pandas as pd
import numpy as np

from sklearn.model_selection import train_test_split

from sklearn.tree import DecisionTreeClassifier

from sklearn.ensemble import RandomForestClassifier

from sklearn.metrics import accuracy_score, classification_report
from google.colab import files
```

Step 2: Create or Upload a Dataset

```
# Check if the user wants to upload a file or generate one
print("Do you have a CSV file to upload? (yes/no)")
response = input().lower()
if response == "yes": # Upload the CSV file
    uploaded = files.upload()
    filename = list(uploaded.keys())[0]
else:
    # Generate synthetic classification data
    from sklearn.datasets import make_classification

    X, y = make_classification(n_samples=300, n_features=10, n_classes=2, random_state=42)

    # Combine features and target into a DataFrame
    data = pd.DataFrame(X, columns=[f'Feature_{i}' for i in range(X.shape[1])])
    data['Target'] = y

    # Save the synthetic dataset to a CSV file
    filename = "synthetic_data.csv"
    data.to_csv(filename, index=False)
    print(f'Synthetic dataset saved as {filename}.')
```

Step 3: Load the Dataset

Load the dataset

```
data = pd.read_csv(filename)

# Display the first few rows of the dataset
print("Dataset Preview:")
print(data.head())
```

Step 4: Preprocess the Data

Separate features (X) and target (y)

```
X = data.iloc[:, :-1].values # All columns except the last one
```

```
y = data.iloc[:, -1].values # Last column as the target
```

Split the dataset into training (80%) and testing (20%) sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Step 5: Train a Single Decision Tree Classifier # Initialize and train the Decision Tree model

```
decision_tree = DecisionTreeClassifier(random_state=42)
```

```
decision_tree.fit(X_train, y_train) #
```

Predict and evaluate

```
y_pred_tree = decision_tree.predict(X_test)
```

```
accuracy_tree = accuracy_score(y_test, y_pred_tree)
```

```
print(f"Decision Tree Accuracy: {accuracy_tree:.2f}")
```

Step 6: Train a Random Forest Classifier

Initialize the Random Forest model with hyperparameter tuning

```
random_forest = RandomForestClassifier(n_estimators=100, max_features='sqrt',
```

```
random_state=42) # Train the model
```

```
random_forest.fit(X_train, y_train)
```

Predict and evaluate

```
y_pred_rf = random_forest.predict(X_test) accuracy_rf
```

```
= accuracy_score(y_test, y_pred_rf)
```

```
print(f"Random Forest Accuracy (100 trees, sqrt features): {accuracy_rf:.2f}")
```

Step 7: Experiment with Random Forest Hyperparameters # Experiment with fewer trees and different feature sampling

```
rf_experiment = RandomForestClassifier(n_estimators=50, max_features=3,  
random_state=42)
```

```
rf_experiment.fit(X_train, y_train)
```

Predict and evaluate

```
y_pred_rf_exp = rf_experiment.predict(X_test) accuracy_rf_exp
```

```
= accuracy_score(y_test, y_pred_rf_exp)
```

```
print(f"Random Forest Accuracy (50 trees, max_features=3): {accuracy_rf_exp:.2f}")
```

Step 8: Compare the Models

print("\nModel Comparison:")

print(f'Decision Tree Accuracy: {accuracy_tree:.2f}') print(f'Random

Forest Accuracy (100 trees): {accuracy_rf:.2f}')

print(f'Random Forest Accuracy (50 trees, max_features=3): {accuracy_rf_exp:.2f}')

Step 9: Visualize Feature Importance (Optional)

import matplotlib.pyplot as plt

Extract feature importance from the Random Forest model

feature_importances = random_forest.feature_importances_

Plot the feature importance plt.figure(figsize=(10,


6))

plt.bar(range(len(feature_importances)), feature_importances, tick_label=data.columns[:-1])

plt.title("Feature Importance in Random Forest") plt.xlabel("Features")

plt.ylabel("Importance") plt.show()

Output :



4f. Implement a gradient boosting machine (e.g., XGBoost). Tune hyperparameters and explore feature importance.

Code :

Step 1: Import Required Libraries

Import necessary libraries

```
import pandas as pd
import numpy as np

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import accuracy_score, classification_report
from xgboost import XGBClassifier, plot_importance
import matplotlib.pyplot as plt
from google.colab import files
```

Step 2: Create or Upload a Dataset

Check if the user wants to upload a file or generate

one print("Do you have a CSV file to upload? (yes/no)")

```
response = input().lower()
if response == "yes": # Upload the CSV file
    uploaded = files.upload()
    filename = list(uploaded.keys())[0]
else:
```

Generate synthetic classification data from

```
sklearn.datasets import make_classification
```

```
X, y = make_classification(n_samples=300, n_features=10, n_classes=2, random_state=42)
```

Combine features and target into a DataFrame

```
data = pd.DataFrame(X, columns=[f'Feature_{i}'
```

```
for i in range(X.shape[1])]
data['Target'] = y #
```

Save the synthetic dataset to a CSV file

```
filename = "synthetic_data.csv"
```

```
data.to_csv(filename, index=False)
```

```
print(f'Synthetic dataset saved as {filename}.')
```

Step 3: Load the Dataset # Load

the dataset data =

```
pd.read_csv(filename) # Display
```

the first few rows of the dataset

```
print("Dataset Preview:")
```

```
print(data.head())
```

Step 4: Preprocess the Data

Separate features (X) and target (y)

```
X = data.iloc[:, :-1].values # All columns except the last one
y = data.iloc[:, -1].values # Last column as the target
```

Split the dataset into training (80%) and testing (20%) sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Step 5: Train a Basic XGBoost Model

Initialize and train the XGBoost model with default parameters

```
xgb = XGBClassifier(random_state=42) xgb.fit(X_train, y_train)
```

Predict and evaluate the model y_pred

```
= xgb.predict(X_test)
```

```
accuracy = accuracy_score(y_test, y_pred)
```

```
print(f'XGBoost Accuracy (Default Parameters): {accuracy:.2f}') Step
```

6: Tune Hyperparameters with GridSearchCV

Define a grid of hyperparameters

```
param_grid = { 'n_estimators': [50, 100, 150], 'learning_rate': [0.01, 0.1, 0.2], 'max_depth': [3, 5, 7] }
```

Initialize GridSearchCV

```
grid_search =
```

```
GridSearchCV(estimator=XGBClassifier(random_state=42), param_grid=param_grid, scoring='accuracy', cv=3, verbose=1) # Fit the model with grid search
```

```
grid_search.fit(X_train, y_train)
```

Best parameters from GridSearch

```
print(f'Best Parameters: {grid_search.best_params_}') # Train the final model
```

```
with the best parameters best_xgb = grid_search.best_estimator_
```

```
# Predict and evaluate y_pred_best =
```

```
best_xgb.predict(X_test) accuracy_best =
```

```
accuracy_score(y_test, y_pred_best)
```

```
print(f'XGBoost Accuracy (Tuned Parameters): {accuracy_best:.2f}')
```

Step 7: Explore Feature Importance

Plot feature importance for the tuned model plt.figure(figsize=(10, 6))


```
plot_importance(best_xgb, importance_type='weight', xlabel="Importance", ylabel="Features") plt.title("XGBoost Feature Importance") plt.show()
```


Step 8: Evaluate the Model

Print a detailed classification report

```
print("Classification Report:")  
print(classification_report(y_test,  
y_pred_best))
```

Output :



Practical 5

5a. Implement and demonstrate the working of a Naive Bayesian classifier using a sample data set. Build the model to classify a test sample.

Step 1: Import Required Libraries # Import necessary

```
libraries import pandas as pd import numpy as np from  
sklearn.model_selection import train_test_split from  
sklearn.metrics import accuracy_score, classification_report from  
sklearn.naive_bayes import GaussianNB  
from google.colab import files
```

Step 2: Create or Upload a Dataset

Ask if the user wants to upload a file or generate one

```
print("Do you have a CSV file to upload? (yes/no)")
```

```

response = input().lower() if response == "yes": # Upload
the CSV file uploaded =files.upload() filename =
list(uploaded.keys())[0]
else:

    # Generate synthetic classification data from
sklearn.datasets import make_classification

    X, y = make_classification(n_samples=300, n_features=8, n_classes=2, random_state=42)

    # Combine features and target into a DataFrame

data = pd.DataFrame(X, columns=[f'Feature_{i}' for i in range(X.shape[1])])
data['Target'] = y
    # Save the synthetic dataset to a CSV file

filename="synthetic_naive_bayes_data.csv"
data.to_csv(filename,index=False)
print(f'Synthetic dataset saved as
{filename}.') Step 3: Load the Dataset #
Load the dataset data =pd.read_csv(filename)

# Display the first few rows of the dataset
print("Dataset Preview:") print(data.head())

```

Step 4: Preprocess the Data

Separate features (X) and target (y)

```

X = data.iloc[:, :-1].values # All columns except the last one
y = data.iloc[:, -1].values # Last column as the target

```

Split the dataset into training (80%) and testing (20%) sets

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

```

Step 5: Train a Naive Bayes Classifier

Initialize the Gaussian Naive Bayes

```

classifier naive_bayes = GaussianNB() # Train
the model

```

```

naive_bayes.fit(X_train, y_train)

```

Step 6: Make Predictions and Evaluate

Predict on the test set

```
y_pred = naive_bayes.predict(X_test)
```

Evaluate the model

```
accuracy = accuracy_score(y_test, y_pred)
print(f'Naive Bayes Accuracy:
{accuracy:.2f}')
```

Detailed classification report

```
print("Classification Report:")
print(classification_report(y_test,
y_pred))
```

Step 7: Test the Model with a Custom Sample

```
# Define a sample test input (replace with meaningful values based on your dataset) test_sample
= [X_test[0]]
```

Taking the first test sample for demonstration

Predict the class for the test sample

```
predicted_class =
naive_bayes.predict(test_sample) print(f'Test
Sample: {test_sample}') print(f'Predicted Class:
{predicted_class[0]}')
```

Output :



	Feature_0	Feature_1	Feature_2	Feature_3	Feature_4	Feature_5	\
0	-1.274158	1.317988	-2.423879	0.906946	-1.583903	-0.331811	
1	1.607963	-1.649959	0.299293	-0.891720	1.301741	1.508502	
2	-0.154167	0.161033	2.210523	0.139400	-0.557492	0.087713	
3	-0.920991	0.949136	-1.613561	0.588410	1.471170	-0.529287	
4	1.013304	-1.038578	-0.305225	-0.539334	-0.609512	1.048078	

	Feature_6	Feature_7	Target
0	-0.452306	0.760415	1
1	0.742095	1.561511	0
2	0.963879	-1.369803	0
3	-1.371901	-0.209324	0
4	-1.065114	-0.186971	0

```
➦ Test Sample: [array([-0.90320608,  0.9220511 , -1.32308979,  0.41081065,  1.64201516,
-1.23559176, -0.63896175,  1.00981709])]
Predicted Class: 1
```

5b. Implement Hidden Markov Models using hmmlearn

Code :

Step 1: Install Required Libraries

Install hmmlearn

```
!pip install hmmlearn
```

Step 2: Import Required Libraries

Import necessary libraries

```
import numpy as np import pandas
```

```
as pd from hmmlearn import hmm
```

```
import matplotlib.pyplot as plt
```

Step 3: Create or Load a Dataset

```
# Generate synthetic observable data np.random.seed(42)
```

```
# Create a sequence of observations and hidden states
```

```
observations = np.random.choice(['A', 'B', 'C'], size=100, p=[0.5, 0.3, 0.2])
```

```
hidden_states = np.random.choice(['X', 'Y'], size=100, p=[0.6, 0.4])
```

```
# Save the data in a DataFrame for analysis data =
```

```
pd.DataFrame({'Observations': observations, 'Hidden States': hidden_states})
```

```
print("Generated Data:")
```

```
print(data.head())
```

Step 4: Encode Observations

```
# Encode the observations into integers observation_mapping = {obs: idx for idx, obs
in enumerate(np.unique(observations))} encoded_observations =
np.array([observation_mapping[obs] for obs in observations])
```

```
# Print the mapping print("Observation
```

```
Encoding:") print(observation_mapping)
```

Step 5: Initialize and Configure the HMM

```
# Initialize the HMM model n_states = 2 # Number of hidden states n_observations =
len(observation_mapping) # Number of unique observations model =
hmm.MultinomialHMM(n_components=n_states, random_state=42, n_iter=100,
tol=0.01)

# Define start probabilities (initial distribution of states)
start_probs = np.array([0.6, 0.4]) # Assumed probabilities
model.startprob_ = start_probs

# Define transition probabilities between states trans_probs
= np.array([
    [0.7, 0.3], # From state X
    [0.4, 0.6], # From state Y]) model.transmat_ =
trans_probs

# Define emission probabilities (probability of observations given states) emission_probs
= np.array([
    [0.5, 0.4, 0.1], # State X emits A, B, C
    [0.2, 0.3, 0.5], # State Y emits A, B, C
])
model.emissionprob_ = emission_probs

# Print the configured model parameters
print("Start Probabilities:",
model.startprob_) print("Transition Matrix:",
model.transmat_) print("Emission
Probabilities:", model.emissionprob_)
```

Step 6: Train the Model

```
# Reshape the data for HMM (requires 2D array) encoded_observations
= encoded_observations.reshape(-1, 1)

# Fit the model
model.fit(encoded_observations)

# Predict hidden states for the observations predicted_states
= model.predict(encoded_observations) # Print the predicted
states print("Predicted States:")
print(predicted_states)
```

Step 7: Visualize the Results

Map predicted states back to their original labels

```
state_mapping = {0: 'X', 1: 'Y'}
```

```
predicted_state_labels = [state_mapping[state] for state in predicted_states]
```

Add predicted states to the DataFrame data['Predicted States'] = predicted_state_labels

Display the first few rows with predicted states

```
print("Data with Predicted States:")
```

```
print(data.head())
```

Plot the observations and predicted states

```
plt.figure(figsize=(12, 6))
```

```
plt.plot(data['Observations'], label='Observations', marker='o', linestyle='-', alpha=0.7)
```

```
plt.plot(data['Predicted States'], label='Predicted States', marker='x', linestyle='--', alpha=0.7)
```


```
plt.legend()
```

```
plt.title("Observations and Predicted
```

```
States") plt.xlabel("Time")
```

```
plt.ylabel("Value") plt.show()
```

Output :



Practical 6 : Probabilistic Model

6a. Implement Bayesian Linear Regression to explore prior and posterior distribution.

Bayesian Linear Regression is a probabilistic approach to linear regression that incorporates uncertainty in the model parameters. Instead of estimating point

values for parameters (as in traditional linear regression), we estimate distributions over the parameters.

Code :

Step 1: Install Required Libraries

Install necessary libraries

```
!pip install matplotlib seaborn scikit-learn
```

Step 2: Import Required Libraries

Import necessary libraries

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.linear_model import BayesianRidge
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from google.colab import files
```

Step 3: Create or Upload a Dataset #

Upload a CSV file if you have one

```
print("Do you have a CSV file to upload? (yes/no)")
response = input().lower()
if response == "yes":
```

Upload the CSV file

```
    uploaded = files.upload()
```

```
    filename =
```

```
    list(uploaded.keys())[0]
```

```
else:
```

Generate synthetic data for demonstration

```
np.random.seed(42)
```

```
    X = np.random.rand(100, 1) * 10 #
```

Random data between 0 and 10 $y = 2 * X$

```
+ 1 + np.random.randn(100, 1) * 2
```

$y = 2x + 1$ with some noise

Convert to a DataFrame

```
data = pd.DataFrame(np.hstack((X, y)), columns=["X", "y"])  
# Save to CSV for convenience  
  
filename="synthetic_data.csv"  
  
data.to_csv(filename,index=False)  
  
print(f'Synthetic dataset saved as {filename}.')
```

Step 4: Load and Explore the Data

Load the dataset (for CSV file)

```
data = pd.read_csv(filename) #  
Display first few rows  
print("Dataset Preview:")  
print(data.head())
```

Step 5: Preprocess the Data

Separate features (X) and target (y)

```
X = data["X"].values.reshape(-1, 1) # Feature matrix y  
= data["y"].values # Target vector
```

Split the dataset into training (80%) and testing (20%) sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Step 6: Implement Bayesian Linear Regression Model

Initialize the BayesianRidge model (which implements Bayesian Linear Regression)

```
bayesian_regressor = BayesianRidge() # Fit the model on the training data  
bayesian_regressor.fit(X_train, y_train)
```

Predict on the test data y_pred =

```
bayesian_regressor.predict(X_test)
```

Step 7: Visualize the Prior and Posterior Distributions # Plot

the prior and posterior distributions of the parameters fig,

```
ax = plt.subplots(1, 2, figsize=(12, 6))
```



```

# Plot prior distribution (assuming the model starts with a standard prior)
ax[0].set_title("Prior Distribution (Assumed)") ax[0].hist(np.random.normal(0, 1, 1000),
bins=50, alpha=0.7, color='blue', label="Prior") ax[0].legend()

# Plot posterior distribution (after model fitting) ax[1].set_title("Posterior
Distribution (After Fitting)")

ax[1].hist(bayesian_regressor.coef_, bins=50, alpha=0.7, color='green', label="Posterior")
ax[1].legend()

plt.show()

```

Step 8: Evaluate the Model Performance #

Calculate the Mean Squared Error (MSE)

```

mse = mean_squared_error(y_test, y_pred)
print(f"Mean Squared Error (MSE):
{mse:.2f}")

```

Step 9: Visualize the Fit of the Model

Plot the true values and the predicted values

```



plt.figure(figsize=(8, 6))

plt.scatter(X_test, y_test, color="blue", label="True values")
plt.plot(X_test, y_pred, color="red", label="Predicted values",
linewidth=2) plt.title("Bayesian Linear Regression: True vs Predicted
Values") plt.xlabel("X") plt.ylabel("y") plt.legend() plt.show()

```

Mean Squared Error (MSE): 3.9

Output :



6b. Implement Gaussian Mixture Models for density estimation and unsupervised clustering.

Code :

Step 1: Install Required Libraries

Install required libraries

```
!pip install matplotlib seaborn scikit-learn
```

Step 2: Import Required Libraries

Import necessary libraries

```
import numpy as np import pandas
```

```
as pd import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
from sklearn.mixture import GaussianMixture from  
sklearn.model_selection import train_test_split from  
google.colab import files
```

Step 3: Create or Upload a Dataset

```
#Ask if the user has a CSV file to upload print("Do  
you have a CSV file to upload? (yes/no)") response  
= input().lower() if response == "yes": # Upload the  
CSV file uploaded = files.upload() filename =  
list(uploaded.keys())[0] else:
```

```
    # Generate synthetic 2D data with two clusters for demonstration np.random.seed(42)  
    # Generate data for two Gaussian distributions
```

```
    X1 = np.random.normal(loc=0, scale=1, size=(300, 2)) # Cluster 1: mean = 0, std = 1
```

```
    X2 = np.random.normal(loc=5, scale=1, size=(300, 2)) # Cluster 2: mean = 5, std = 1    #
```

Stack the data to create a dataset

```
    X = np.vstack([X1, X2])
```

Create DataFrame to simulate the CSV file for consistency

```
data = pd.DataFrame(X, columns=["Feature_1", "Feature_2"])  
filename = "synthetic_data.csv" data.to_csv(filename,  
index=False) print(f"Synthetic dataset saved as {filename}.")
```

Step 4: Load and Explore the Dataset

Load the dataset (if CSV file is uploaded)

```
data = pd.read_csv(filename) # Display the  
first few rows print("Dataset Preview:")  
print(data.head())
```

Plot the data to visualize its structure

```
sns.scatterplot(data=data, x="Feature_1",  
y="Feature_2") plt.title("Synthetic Data")  
plt.xlabel("Feature 1") plt.ylabel("Feature 2") plt.show()
```

Step 5: Fit a Gaussian Mixture Model (GMM)

Define the GMM model

```
n_components = 2 # Number of Gaussian distributions (clusters) gmm =  
GaussianMixture(n_components=n_components, covariance_type='full',  
random_state=42)
```

Fit the GMM model to the data gmm.fit(data)

Predict the cluster labels for each data point labels

```
= gmm.predict(data)
```

Add the cluster labels to the dataset for visualization

```
data['Cluster'] = labels # Plot the clustered data
```

```
sns.scatterplot(data=data, x="Feature_1", y="Feature_2", hue="Cluster", palette="viridis",  
marker="o")
```

```
plt.title("Gaussian Mixture Model Clustering")
```

```
plt.xlabel("Feature 1") plt.ylabel("Feature 2")
```

```
plt.legend()
```

```
plt.show()
```

Step 6: Visualize the Gaussian Mixture Model (GMM) Components

Extract the means and covariances of the Gaussian components

```
means = gmm.means_ covariances = gmm.covariances_
```

```
# Plot the GMM components on top of the data plt.figure(figsize=(8,  
6)) # Plot data points sns.scatterplot(data=data, x="Feature_1",  
y="Feature_2", hue="Cluster", palette="viridis", marker="o", s=60,  
alpha=0.7)
```

Plot the GMM ellipses for mean, covar in zip(means, covariances):

Plot the Gaussian components as ellipses

```
v, w = np.linalg.eigh(covar) v = 2.0 *  
np.sqrt(2.0) * np.sqrt(v)
```

```
# Scaling factor for the ellipse
```

```
u = w[0] / np.linalg.norm(w[0])
```

```
# Normalize the eigenvector
```

```
angle = np.arctan(u[1] / u[0])
```

```
# Create the ellipse angle = angle * 180.0 / np.pi #
```

```
Convert to degrees
```

```
ellipse = plt.matplotlib.patches.Ellipse(mean, v[0], v[1], angle=angle, color='red', alpha=0.3)  
plt.gca().add_patch(ellipse)
```

```
plt.title("GMM Clustering with Gaussian Components")
```

```
plt.xlabel("Feature 1") plt.ylabel("Feature 2")
```

```
plt.legend()
```

```
plt.show()
```

Step 7: Model Evaluation (Optional)

```
# Compute the log-likelihood of the data under the fitted GMM model log_likelihood  
= gmm.score(data)
```

```
print(f'Log-Likelihood of the data: {log_likelihood:.2f}')
```

Step 8: Predict New Data Points

```
# Example of predicting the cluster for new data points
```

```
new_data = np.array([[1.5, 2.5], [4.5, 5.5], [7.0, 8.0]])
```

```
new_labels = gmm.predict(new_data)
```



```
# Print the predicted clusters for the new data
```


```
points print("Predicted Clusters for New Data
```

```
Points:") for i, label in enumerate(new_labels):
```

```
    print(f'Data point {new_data[i]} is in Cluster {label}')
```

Output :





Practical 7 : Model Evaluation and Hyperparameter Tuning

7a. Implement cross-validation techniques (k-fold, stratified, etc.) for robust model evaluation

Code :

1. Import Necessary Libraries

```
import numpy as np
import pandas as pd

from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split, KFold, StratifiedKFold, GridSearchCV
from sklearn.ensemble import RandomForestClassifier

from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

import matplotlib.pyplot as plt
import seaborn as sns
```

2. Generate a Synthetic Dataset

```
# Create a synthetic dataset with 2 classes
```

```

X, y = make_classification( n_samples=1000, n_features=10,
    n_informative=8,    n_redundant=2,    n_clusters_per_class=1,
    random_state=42
)
# Convert to a DataFrame for visualization
df = pd.DataFrame(X, columns=[f'Feature_{i}' for i in range(1, 11)]) df['Target']
    = y
# Display the first few rows
print(df.head())

```

3. Split Data into Train and Test Sets

Split data into 80% training and 20% testing

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42,
    stratify=y)

```

```

4. Define k-Fold Cross-Validation kf = KFold(n_splits=5,
    shuffle=True,    random_state=42) print("k-Fold Cross-
    Validation:") for train_index, val_index in kf.split(X_train):
    print("TRAIN:", train_index, "VALIDATION:", val_index)

```

```

5. Define Stratified k-Fold Cross-Validation skf =
    StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
    print("\nStratified k-Fold Cross-Validation:") for train_index,
    val_index in skf.split(X_train, y_train): print("TRAIN:",
    train_index, "VALIDATION:", val_index)

```

6. Train and Evaluate Using k-Fold Cross-Validation

Initialize model

```

model = RandomForestClassifier(random_state=42)

```


Perform k-Fold Cross-Validation

```
accuracies = [] for train_index, val_index in
kf.split(X_train):
    X_kf_train, X_kf_val = X_train[train_index], X_train[val_index] y_kf_train,
    y_kf_val = y_train[train_index], y_train[val_index]
# Train model
    model.fit(X_kf_train, y_kf_train)
# Validate model
    y_pred = model.predict(X_kf_val) accuracy =
accuracy_score(y_kf_val, y_pred) accuracies.append(accuracy)
print(f'Average Accuracy from k-Fold: {np.mean(accuracies):.2f}')
```

7. Hyperparameter Tuning Using GridSearchCV

```
# Define parameter grid param_grid = {
    'n_estimators': [50, 100, 200],
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [2, 5, 10],
}
# Perform GridSearchCV with Stratified k-Fold grid_search
= GridSearchCV(
    estimator=RandomForestClassifier(random_state=42),
    param_grid=param_grid,
    cv=StratifiedKFold(n_splits=5, shuffle=True,
    random_state=42), scoring='accuracy', n_jobs=-1, verbose=1
)
# Fit to training data grid_search.fit(X_train,
y_train)
print("Best Parameters:", grid_search.best_params_)
print("Best Cross-Validation Accuracy:", grid_search.best_score_)
```

8. Evaluate the Final Model # Use the

best model for evaluation best_model =

grid_search.best_estimator_

Predict on test data

y_test_pred = best_model.predict(X_test)

Evaluate performance print("\nTest Accuracy:", accuracy_score(y_test,

y_test_pred)) print("\nClassification Report:\n", classification_report(y_test,

y_test_pred)) **# Confusion matrix**

conf_matrix = confusion_matrix(y_test, y_test_pred)

sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=['Class 0', 'Class 1'], yticklabels=['Class 0', 'Class 1'])

plt.xlabel('Predicted')

plt.ylabel('Actual')

plt.title('Confusion Matrix')

plt.show()

Output :



```
Feature_1 Feature_2 Feature_3 Feature_4 Feature_5 Feature_6 \
0 -3.358483 3.159918 0.827163 0.069638 -6.715639 -2.708559
1 2.071819 -4.055419 -2.615940 -2.599432 3.053752 0.366795
2 -0.633460 0.712482 2.024390 -0.432639 -1.307929 0.419320
3 -0.464478 0.892442 2.521010 2.766580 -1.933734 -1.418018
4 1.042426 -1.192605 -2.071386 -0.131231 0.545377 0.379060

Feature_7 Feature_8 Feature_9 Feature_10 Target
0 0.183206 1.113502 1.730759 1.228394 1
1 -0.392171 -1.191720 -1.220516 1.899925 0
2 -1.469510 -0.719051 1.155005 2.018026 0
3 1.391760 -2.430279 1.308295 -0.270896 1
4 -0.062978 -1.325591 2.037936 0.115414 0
k-Fold Cross-Validation:
TRAIN: [ 0 1 3 4 5 6 8 9 11 12 13 14 15 16 17 18 19 20
21 22 24 25 26 27 28 32 34 35 36 37 38 40 41 42 43 44
45 46 47 48 50 51 52 53 55 56 57 58 59 60 61 62 64 68
69 70 71 73 74 75 79 80 82 83 85 87 88 89 90 91 92 93
94 95 98 99 100 102 103 104 105 106 107 108 111 112 113 114 115 116
117 119 121 122 123 124 125 126 127 128 129 130 131 132 133 134 135 136
138 140 141 142 143 144 145 146 147 148 149 150 151 152 153 154 156 157
158 159 160 161 162 163 164 165 166 167 169 170 171 172 173 175 176 177
178 179 180 181 182 183 184 185 186 187 188 189 190 191 193 194 195 196
197 200 201 202 203 205 206 207 212 213 214 216 217 219 220 221 222 223
224 225 226 227 228 229 230 232 233 234 236 237 238 239 240 241 242 243
245 246 247 248 249 251 252 253 255 256 257 258 259 261 262 263 264 267
268 269 270 271 272 273 274 276 277 278 279 280 282 283 284 285 287 288
289 290 291 292 293 295 297 298 299 300 301 303 304 305 307 308 309 310
311 312 313 315 317 318 319 320 321 322 324 325 328 329 330 331 332 334
```




```
674 675 676 677 678 679 680 681 682 683 684 685 686 687 689 690 691 693
694 695 697 698 699 700 701 702 703 704 706 707 708 709 710 711 712 713
714 716 717 718 722 724 725 726 727 728 729 730 731 732 733 734 735 736
737 739 740 742 743 745 747 748 749 751 752 753 755 756 757 759 761 762
763 765 766 767 768 769 770 771 772 773 774 775 778 779 780 781 782 784
785 788 789 790 791 792 793 794 797 799] VALIDATION: [ 2 7 10 23 29 30 31 33 39 49 54 63 65 66 67 72 76 77
78 81 84 86 96 97 101 109 110 118 120 137 139 155 168 174 192 198
199 204 208 209 210 211 215 218 231 235 244 250 254 260 265 266 275 281
286 294 296 302 306 314 316 323 326 327 333 336 346 352 357 360 361 365
367 368 377 383 388 393 395 398 409 422 423 425 428 432 446 456 464 481
483 486 490 506 512 513 515 519 521 525 526 529 532 533 534 537 545 568
570 578 589 591 594 595 596 604 608 610 621 622 628 635 637 640 641 644
652 655 656 658 662 666 667 670 688 692 696 705 715 719 720 721 723 738
741 744 746 750 754 758 760 764 776 777 783 786 787 795 796 798]
TRAIN: [ 0 1 2 3 4 5 7 8 9 10 12 13 14 15 16 17 19 20
21 22 23 25 26 27 29 30 31 32 33 34 35 36 37 38 39 40
45 46 47 48 49 50 52 53 54 57 58 59 62 63 64 65 66 67
68 71 72 75 76 77 78 80 81 84 85 86 87 88 91 93 94 95
96 97 98 99 100 101 102 103 105 106 107 109 110 111 112 113 115 116
117 118 119 120 121 122 123 124 125 126 127 128 129 130 134 137 138 139
141 142 143 144 146 147 149 150 151 152 153 154 155 156 157 159 160 161
162 166 168 169 170 171 172 173 174 175 179 180 183 184 185 186 187 188
189 190 191 192 194 195 197 198 199 200 201 202 203 204 205 206 207 208
209 210 211 214 215 216 217 218 219 221 222 224 225 226 228 229 230 231
232 233 235 236 237 238 240 241 242 243 244 245 246 249 250 251 252 253
254 255 256 257 258 260 261 262 263 265 266 267 268 269 270 271 272 273
274 275 276 277 278 279 280 281 282 283 284 286 287 288 289 293 294 295
296 297 298 301 302 303 304 305 306 307 308 310 312 313 314 315 316 317
318 320 321 322 323 324 325 326 327 330 333 335 336 337 338 339 340 341
```

```
450 455 459 460 463 470 472 475 480 481 483 494 496 502 503 505 514 519
521 523 529 533 553 555 560 563 565 579 585 590 594 600 603 620 630 631
634 635 637 638 644 646 648 649 651 673 675 686 691 693 708 709 715 720
729 730 738 747 753 759 767 771 774 776 777 780 782 783 784 796]
Average Accuracy from k-Fold: 0.95
Fitting 5 folds for each of 36 candidates, totalling 180 fits
Best Parameters: {'max_depth': None, 'min_samples_split': 5, 'n_estimators': 50}
Best Cross-Validation Accuracy: 0.96
```

Test Accuracy: 0.955

Classification Report:

	precision	recall	f1-score	support
0	1.00	0.91	0.95	100
1	0.92	1.00	0.96	100
accuracy			0.95	200
macro avg	0.96	0.96	0.95	200
weighted avg	0.96	0.95	0.95	200



7b. Systematically explore combinations of hyperparameters to optimize model performance.(use grid and randomized search)

Code :

1. Import Necessary Libraries

```
import numpy as np
import pandas as pd

from sklearn.datasets import make_classification

from sklearn.model_selection import train_test_split, GridSearchCV, RandomizedSearchCV, StratifiedKFold

from sklearn.ensemble import RandomForestClassifier

from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

import matplotlib.pyplot as plt

import seaborn as sns
```

2. Generate a Synthetic Dataset

Generate a binary classification dataset

```
X, y = make_classification( n_samples=1000, n_features=12,
                           n_informative=8, n_redundant=2, n_clusters_per_class=1,
                           flip_y=0.03, random_state=42
)
```

Convert to a DataFrame for visualization

```
df = pd.DataFrame(X, columns=[f'Feature_{i}' for i in range(1, 13)])
df['Target'] = y
```

Display the first few rows

```
print(df.head())
```

3. Split Data into Train and Test Sets

Split data into training and testing sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)
```

4. Define the Model

Initialize a Random Forest classifier

```
model = RandomForestClassifier(random_state=42)
```

5. Hyperparameter Tuning Using Grid Search # Define a

```
parameter grid for Grid Search param_grid = {  
    'n_estimators': [50, 100, 200],  
    'max_depth': [None, 10, 20],  
    'min_samples_split': [2, 5, 10],  
    'min_samples_leaf': [1, 2, 4]  
}
```

GridSearchCV with 5-fold cross-validation

```
grid_search = GridSearchCV(  
    estimator=model, param_grid=param_grid,  
    scoring='accuracy',  
    cv=StratifiedKFold(n_splits=5, shuffle=True,  
        random_state=42), verbose=1, n_jobs=-1  
)
```

Fit the model

```
grid_search.fit(X_train, y_train)
```

Best parameters and score from Grid Search

```
print("Best Parameters from Grid Search:", grid_search.best_params_)  
print("Best Cross-Validation Accuracy from Grid Search:", grid_search.best_score_)
```

6. Hyperparameter Tuning Using Randomized Search

```
from scipy.stats import randint
```

Define a parameter distribution for Randomized Search

```
param_dist = {  
    'n_estimators': randint(50, 300),  
    'max_depth': [None, 10, 20, 30],  
    'min_samples_split': randint(2, 15),  
    'min_samples_leaf': randint(1, 10)  
}
```

RandomizedSearchCV with 5-fold cross-validation

```
random_search = RandomizedSearchCV(
    estimator=model,
    param_distributions=param_dist,
    n_iter=50, # Number of random combinations to try scoring='accuracy',
    cv=StratifiedKFold(n_splits=5, shuffle=True,
    random_state=42), verbose=1, n_jobs=-1, random_state=42
)
```

Fit the model

```
random_search.fit(X_train, y_train)
```

Best parameters and score from Randomized Search

```
print("Best Parameters from Randomized Search:", random_search.best_params_)
print("Best Cross-Validation Accuracy from Randomized Search:",
    random_search.best_score_)
```

7. Evaluate the Best Model

Select the best model from Grid Search and Randomized Search

```
best_model = random_search.best_estimator_ # Or use grid_search.best_estimator_
```

Predict on test data

```
y_test_pred = best_model.predict(X_test)
```

Evaluate the performance

```
print("\nTest Accuracy:", accuracy_score(y_test, y_test_pred))
```

```
print("\nClassification Report:\n", classification_report(y_test, y_test_pred))
```

Confusion Matrix

```
conf_matrix = confusion_matrix(y_test, y_test_pred)
```

```
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=['Class 0', 'Class 1'], yticklabels=['Class 0', 'Class 1'])
```

```
plt.xlabel('Predicted')
```

```
plt.ylabel('Actual')
```

```
plt.title('Confusion Matrix')
```

```
plt.show()
```

Output :


```
Feature_1 Feature_2 Feature_3 Feature_4 Feature_5 Feature_6 \
0 0.013650 0.607473 -2.096916 2.867232 2.504360 0.784101
1 0.107199 0.105735 -3.843343 1.524052 -1.619824 0.778334
2 -1.779086 -5.219831 -0.738488 2.108084 -0.803833 -3.431122
3 -4.310656 -2.268569 1.864943 -1.246116 1.260794 -2.007664
4 -3.195179 -0.671327 3.720485 0.356661 0.819486 2.670230

Feature_7 Feature_8 Feature_9 Feature_10 Feature_11 Feature_12 Target
0 -0.497744 -0.482072 1.112773 1.641637 -2.689832 -0.480311 1
1 0.551177 -1.843583 -0.110132 -0.494739 -0.985276 -0.978400 0
2 1.346120 -0.858351 -0.792415 -2.260815 0.238780 3.029952 0
3 -0.824133 -2.277449 0.936206 1.255903 1.386278 -0.321200 0
4 1.857477 -3.410944 -1.773719 0.656476 3.534189 -1.704889 0

Fitting 5 folds for each of 81 candidates, totalling 405 fits
Best Parameters from Grid Search: {'max_depth': None, 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 100}
Best Cross-Validation Accuracy from Grid Search: 0.9487499999999999
Fitting 5 folds for each of 50 candidates, totalling 250 fits
Best Parameters from Randomized Search: {'max_depth': None, 'min_samples_leaf': 2, 'min_samples_split': 9, 'n_estimators': 285}
Best Cross-Validation Accuracy from Randomized Search: 0.9487499999999999

Test Accuracy: 0.94
```

Classification Report:					
	precision	recall	f1-score	support	
0	0.99	0.89	0.94	100	
1	0.90	0.99	0.94	100	
accuracy			0.94	200	
macro avg	0.94	0.94	0.94	200	
weighted avg	0.94	0.94	0.94	200	



Practical 8 : Bayesian Learning

Implement Bayesian Learning using inferences

Code :

1. **Import Necessary Libraries** import numpy as np import pandas as pd


```
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt
```

2. Generate a Synthetic Dataset

We create a dataset suitable for classification problems.

Generate a dataset with 2 classes

```
X, y = make_classification( n_samples=1000, n_features=8,
n_informative=6, n_redundant=2, n_classes=2, random_state=42) #
```

Convert to DataFrame for visualization df = pd.DataFrame(X,
columns=[f'Feature_{i}' for i in range(1, 9)]) df['Target'] = y

Display the first few rows

```
print(df.head())
```

3. Split the Dataset

Divide the data into training and testing sets.

Split data into 80% training and 20% testing

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42,
stratify=y)
```

4. Bayesian Learning with Naive Bayes

Here, we implement Bayesian Learning using the Gaussian Naive Bayes classifier.

Initialize the Gaussian Naive Bayes model model

```
= GaussianNB()
```

Fit the model to the training data

```
model.fit(X_train, y_train) # Predict
```

on the test data

```
y_pred = model.predict(X_test)
```

5. Evaluate the Model

We evaluate the model's performance using accuracy, classification report, and confusion matrix.

```
# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"Test Accuracy: {accuracy:.2f}")

# Print classification report
print("\nClassification Report:\n", classification_report(y_test, y_pred))

# Generate and plot confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)

sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=['Class 0', 'Class 1'], yticklabels=['Class 0', 'Class 1'])

plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```

6. Understanding Bayesian Inference

In Bayesian Learning, the model predicts based on the probabilities:

- **Prior Probability ($P(C)$):** The likelihood of each class based on historical data.
- **Likelihood ($P(X | C)$):** The probability of the data given a class.
- **Posterior Probability ($P(C | X)$):** Calculated using Bayes' theorem:
$$P(C|X) = \frac{P(X|C) \cdot P(C)}{P(X)}$$

Example: Compute posterior probabilities for the first test sample

```
sample = X_test[0].reshape(1, -1)
posterior_probs = model.predict_proba(sample)
print(f"Sample Features: {sample}")
print(f"Posterior Probabilities: {posterior_probs}")
print(f"Predicted Class: {model.predict(sample)}")
```


Output :

	Feature_1	Feature_2	Feature_3	Feature_4	Feature_5	Feature_6	\
0	-1.732538	5.260112	-2.952194	-4.603768	2.235848	1.928893	
1	2.072914	2.240572	-1.385104	-2.514962	-0.984756	1.436260	
2	-0.263106	1.527781	-1.872414	-0.028009	1.612809	3.264194	
3	-0.164349	-0.550131	-0.019503	-0.765000	2.273523	2.084217	
4	-1.419423	1.015324	-0.864441	-0.009297	0.385404	0.449093	

	Feature_7	Feature_8	Target
0	-0.101845	3.193487	0
1	-1.255271	2.089872	0
2	-1.296421	1.537870	0
3	-0.321931	0.426253	0
4	-0.029007	-1.902917	1

Test Accuracy: 0.84

Classification Report:				
	precision	recall	f1-score	support
0	0.88	0.80	0.84	99
1	0.82	0.89	0.85	101
accuracy			0.84	200
macro avg	0.85	0.84	0.84	200
weighted avg	0.85	0.84	0.84	200



Sample Features: [[1.94397502 0.33753289 -0.77968371 2.63667244 -1.62323864 1.87859635
-2.20522563 2.09632056]]
Posterior Probabilities: [[0.07067404 0.92932596]]
Predicted Class: [1]